

# Principles of Distributed Systems

inft-3507

Dr. J.Burns

**ADA University**

Autumn 2025

## Section 5: Coordination

*This content is based on the following public resources: <https://www.distributed-systems.net/index.php/books/ds4/>*

# Mutual exclusion

# Process Critical Sections

## Problem

Several processes in a distributed system want exclusive access to some resource - and we want to avoid concurrent access to this resource: The processes all have a region of code where the concurrency takes place. This is called the **Critical Section**

## Critical Section

A **critical section** is a code segment within a process (or thread) that accesses or modifies shared resources (e.g., variables, data structures, files, database or devices) and must be executed atomically — as an indivisible unit—by at most one **one concurrent process or thread** at a time.

# Mutual exclusion: Example

```
do {
    // Non-critical section (local work, safe for concurrency)

    // Entry section (acquire synchronization primitive)
    while (!can_enter_critical_section()); // Busy-wait or block

    // *** CRITICAL SECTION ***
    // Shared resource access/modification (e.g., balance += amount;)
    // Must be atomic w.r.t. other threads

    // Exit section (release synchronization primitive)
    signal_exit_critical_section();

    // Non-critical section
} while (true);
```

# Mutual exclusion

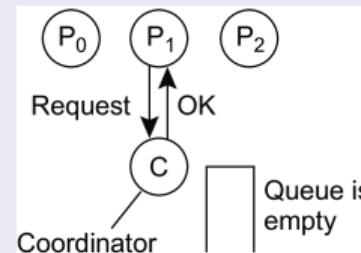
## Basic solutions

**Permission-based:** A process wanting to enter its critical region, or access a resource, needs permission from other processes.

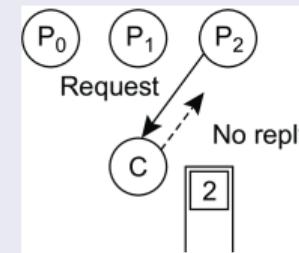
**Token-based:** A token is passed between processes. The one who has the token may proceed in its critical region, or pass it on when not interested.

# Permission-based, centralized

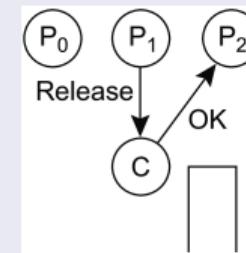
Simply use a coordinator



(a)



(b)



(c)

- (a) Process  $P_1$  asks the coordinator for permission to access a shared resource. Permission is granted.
- (b) Process  $P_2$  then asks permission to access the same resource. The coordinator does not reply.
- (c) When  $P_1$  releases the resource, it tells the coordinator, which then replies to  $P_2$ .

# Mutual exclusion: Ricart & Agrawala

Acknowledgments are not sent (the caller blocks)

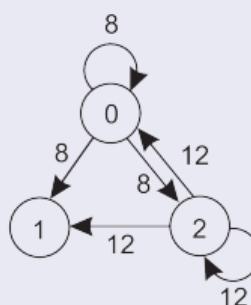
Return a response to a request only when:

- The receiving process has no interest in the shared resource; or
- The receiving process is waiting for the resource, but has lower priority (known through comparison of timestamps).

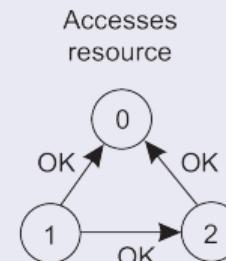
In all other cases, reply is **deferred**, implying some more local administration.

# Mutual exclusion: Ricart & Agrawala

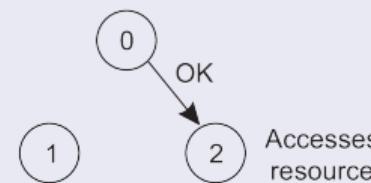
Example with three processes



(a)



(b)



(c)

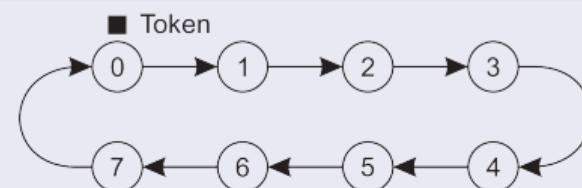
- (a) Two processes want to access a shared resource at the same moment.
- (b)  $P_0$  has the lowest timestamp, so it wins.
- (c) When process  $P_0$  is done, it sends an  $OK$  also, so  $P_2$  can now go ahead.

# Mutual exclusion: Token ring algorithm

## Essence

Organize processes in a **logical** ring, and let a token be passed between them. The one that holds the token is allowed to enter the critical region (if it wants to).

An overlay network constructed as a logical ring with a circulating token



# Decentralized mutual exclusion

## Principle

Assume every resource is replicated  $N$  times, with each replica having its own coordinator  $\Rightarrow$  access requires a **majority vote** from  $m > N/2$  coordinators. A coordinator always responds immediately to a request.

## Assumption

When a coordinator crashes, it will recover quickly, but will have forgotten about permissions it had granted.

# Decentralized mutual exclusion

How robust is this system?

- Let  $p = \Delta t / T$  be the probability that a coordinator resets during a time interval  $\Delta t$ , while having a lifetime of  $T$ .
- The probability  $\mathbb{P}[k]$  that  $k$  out of  $m$  coordinators reset during the same interval is

$$\mathbb{P}[k] = \binom{m}{k} p^k (1-p)^{m-k}$$

- for  $m > N/2$ , if  $f$  coordinators reset  $\Rightarrow$  correctness is violated when there is only a minority of nonfaulty coordinators: when  $N - (m - f) \geq m$ , or,  $f \geq 2m - N$ .
- The probability of a violation is  $\sum_{k=2m-N}^m \mathbb{P}[k]$ .

# Decentralized mutual exclusion

Violation probabilities for various parameter values

N	m	p	Violation
8	5	3 sec/hour	$< 10^{-5}$
8	6	3 sec/hour	$< 10^{-11}$
16	9	3 sec/hour	$< 10^{-4}$
16	12	3 sec/hour	$< 10^{-21}$
32	17	3 sec/hour	$< 10^{-4}$
32	24	3 sec/hour	$< 10^{-43}$

N	m	p	Violation
8	5	30 sec/hour	$< 10^{-3}$
8	6	30 sec/hour	$< 10^{-7}$
16	9	30 sec/hour	$< 10^{-2}$
16	12	30 sec/hour	$< 10^{-13}$
32	17	30 sec/hour	$< 10^{-2}$
32	24	30 sec/hour	$< 10^{-27}$

So

What can we conclude?

# Mutual exclusion: comparison

Algorithm	Messages per entry/exit	Delay before entry (in message times)
Centralized	3	2
Distributed	$2(N - 1)$	$2(N - 1)$
Token ring	$1, \dots, \infty$	$0, \dots, N - 1$
Decentralized	$2kN + (k - 1)N/2 + N, k = 1, 2, \dots$	$2kN + (k - 1)N/2$

Decentralized Case messages per entry/exit for  $k \geq 1$  attempts:

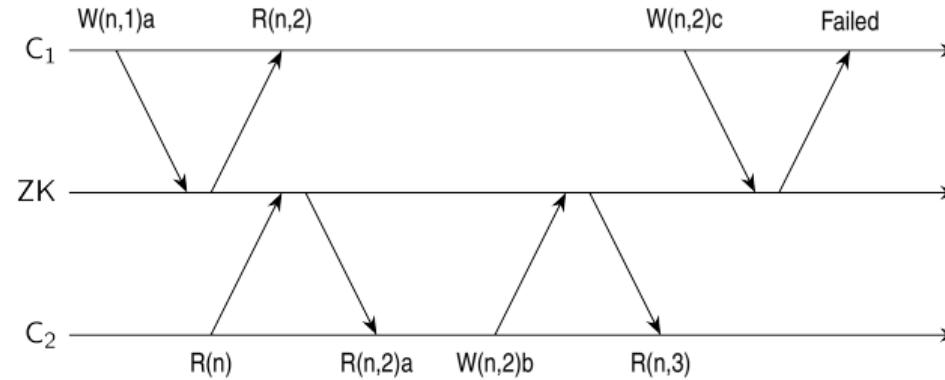
- send  $N$  messages to coordinators receive  $N$  responses.
- If it does not get a majority, release  $N/2$  votes.
- If it did get enough votes, send  $N$  release messages later.
- A process may need to go through  $k \geq 1$  attempts

# Example: ZooKeeper

## Basics (and keeping it simple)

- Centralized server setup
- Distributed Nodes
- All client-server communication is **nonblocking**: a client immediately gets a response
- ZooKeeper maintains a **tree-based namespace**, akin to that of a filesystem
- Clients can **create**, **delete**, or **update** nodes, as well as **check existence**.

## ZooKeeper versioning



## Notations

- $W(n, k)a$ : request to write  $a$  to node  $n$ , assuming current version is  $k$ .
  - $R(n, k)$ : current version of node  $n$  is  $k$ .
  - $R(n)$ : client wants to know the current value of node  $n$
  - $R(n, k)a$ : value  $a$  from node  $n$  is returned with its current version  $k$ .

# ZooKeeper locking protocol

It is now very simple

- ① **lock**: A client  $C_1$  creates a node  $/lock$ .
- ② **lock**: A client  $C_2$  wants to acquire the lock but is notified that the associated node already exists  $\Rightarrow C_2$  subscribes to notification on changes of  $/lock$ .
- ③ **unlock**: Client  $C_1$  deletes node  $/lock \Rightarrow$  all subscribers to changes are notified.

# Election algorithms

# Election algorithms

## Principle

An algorithm requires that some process acts as a coordinator. The question is how to select this special process **dynamically**.

## Note

In many systems, the coordinator is chosen manually (e.g., file servers). This leads to centralized solutions  $\Rightarrow$  single point of failure.

## Teasers

- ① If a coordinator is chosen dynamically, to what extent can we speak about a centralized or distributed solution?
- ② Is a fully distributed solution, i.e. one without a coordinator, always more robust than any centralized/coordinated solution?

# Basic assumptions

- All processes have unique id's
- All processes know id's of all processes in the system (but not if they are up or down)
- Election means identifying the process with the highest id that is up

# Election by bullying

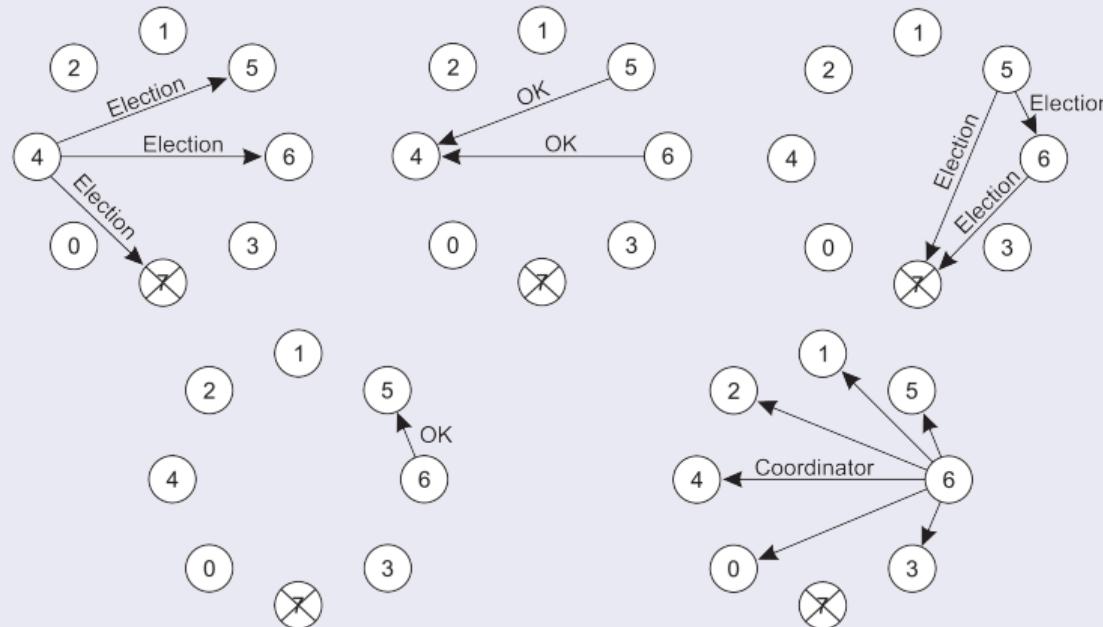
## Principle

Consider  $N$  processes  $\{P_0, \dots, P_{N-1}\}$  and let  $id(P_k) = k$ . When a process  $P_k$  notices that the coordinator is no longer responding to requests, it initiates an election:

- ①  $P_k$  sends an *ELECTION* message to all processes with higher identifiers:  
 $P_{k+1}, P_{k+2}, \dots, P_{N-1}$ .
- ② If no one responds,  $P_k$  wins the election and becomes coordinator.
- ③ If one of the higher-ups answers, it takes over and  $P_k$ 's job is done.

# Election by bullying

## The bully election algorithm



# Election in a ring

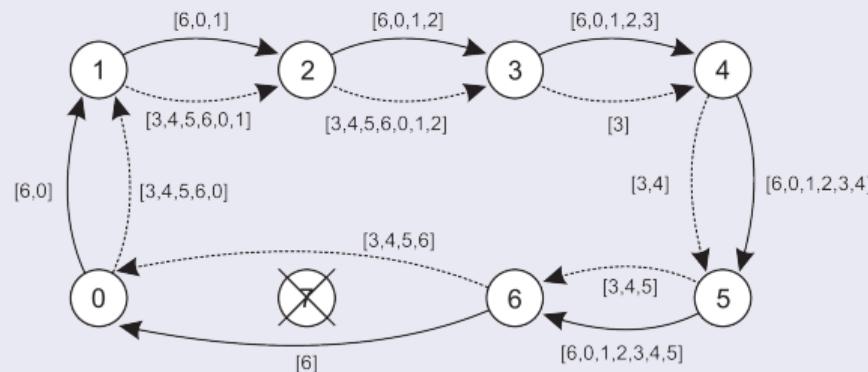
## Principle

Process priority is obtained by organizing processes into a (logical) ring. The process with the highest priority should be elected as coordinator.

- Any process can start an election by sending an election message to its successor. If a successor is down, the message is passed on to the next successor.
- If a message is passed on, the sender adds itself to the list. When it gets back to the initiator, everyone had a chance to make its presence known.
- The initiator sends a coordinator message around the ring containing a list of all living processes. The one with the highest priority is elected as coordinator.

# Election in a ring

## Election algorithm using a ring



- The solid line shows the election messages initiated by  $P_6$
- The dashed one, the messages by  $P_3$

# Example: Leader election **in** ZooKeeper server group

## Basics

- Each server  $s$  in the server group has an identifier  $\text{id}(s)$
- Each server has a monotonically increasing counter  $\text{tx}(s)$  of the latest transaction it handled (i.e., series of operations on the namespace).
- When follower  $s$  suspects leader crashed, it broadcasts an  $ELECTION$  message, along with the pair  $(\text{voteID}, \text{voteTX})$ . Initially,
  - $\text{voteID} \leftarrow \text{id}(s)$
  - $\text{voteTX} \leftarrow \text{tx}(s)$
- Each server  $s$  maintains two variables:
  - $\text{leader}(s)$ : records the server that  $s$  believes may be final leader.  
Initially,  $\text{leader}(s) \leftarrow \text{id}(s)$ .
  - $\text{lastTX}(s)$ : what  $s$  knows to be the most recent transaction.  
Initially,  $\text{lastTX}(s) \leftarrow \text{tx}(s)$ .

## Example: Leader election **in** ZooKeeper server group

When  $s^*$  receives  $(voteID, voteTX)$

- If  $lastTX(s^*) < voteTX$ , then  $s^*$  just received more up-to-date information on the most recent transaction, and sets
  - $leader(s^*) \leftarrow voteID$
  - $lastTX(s^*) \leftarrow voteTX$
- If  $lastTX(s^*) = voteTX$  and  $leader(s^*) < voteID$ , then  $s^*$  knows as much about the most recent transaction as what it was just sent, but its perspective on which server will be the next leader needs to be updated:
  - $leader(s^*) \leftarrow voteID$

### Note

When  $s^*$  believes it should be the leader, it broadcasts  $\langle id(s^*), tx(s^*) \rangle$ . Essentially, we're **bullying**.

# Example: Leader election in Raft

## Basics

- We have a (relatively small) group of servers
- A server is in one of three states: *follower*, *candidate*, or *leader*
- The protocol works in *terms*, starting with term 0
- Each server starts in the *follower* state.
- A leader is to regularly broadcast messages (perhaps just a simple *heartbeat*)

# Example: Leader election in Raft

## Selecting a new leader

When follower  $s^*$  hasn't received anything from the alleged leader  $s$  for some time,  $s^*$  broadcasts that it volunteers to be the next leader, increasing the term by 1.  $s^*$  enters the **candidate** state. Then:

- If leader  $s$  receives the message, it responds by acknowledging that it is still the leader.  $s^*$  returns to the **follower** state.
- If another follower  $s^{**}$  gets the election message from  $s^*$ , and it is the first election message during the current term,  $s^{**}$  votes for  $s^*$ . Otherwise, it simply ignores the election message from  $s^*$ . When  $s^*$  has collected a majority of votes, a new term starts with a new leader.

## Observation

By slightly differing the timeout values per follower for deciding when to start an election, we can avoid concurrent elections, and the election will rapidly converge.

# Elections by proof of work

## Basics

- Consider a potentially large group of processes
- Each process is required to solve a computational puzzle
- When a process solves the puzzle, it broadcasts its victory to the group
- We assume there is a conflict resolution procedure when more than one process claims victory

## Solving a computational puzzle

- Make use of a **secure hashing function**  $H(m)$ :
  - $m$  is some data;  $H(m)$  returns a **fixed-length bit string**
  - computing  $h = H(m)$  is computationally efficient
  - finding a function  $H^{-1}$  such that  $m = H^{-1}(H(m))$  is computationally extremely difficult
- Practice: finding  $H^{-1}$  boils down to an extensive **trial-and-error** procedure

# Elections by proof of work

## Controlled race

- Assume a globally known secure hash function  $H^*$ . Let  $H_i$  be the hash function used by process  $P_i$ .
- Task: given a bit string  $h = H_i(m)$ , find a bit string  $\tilde{h}$  such that  $h^* = H^*(H_i(\tilde{h} \odot h))$  where:
  - $h^*$  is a bit string with  $K$  leading zeroes
  - $\tilde{h} \odot h$  denotes some predetermined bitwise operation on  $\tilde{h}$  and  $h$

## Observation

By controlling  $K$ , we control the difficulty of finding  $\tilde{h}$ . If  $p$  is the probability that a random guess for  $\tilde{h}$  will suffice:  $p = (1/2)^K$ .

## Current practice

In many PoW-based blockchain systems,  $K = 64$

- With  $K = 64$ , it takes about 10 minutes on a supercomputer to find  $\tilde{h}$
- With  $K = 64$ , it takes about 100 years on a laptop to find  $\tilde{h}$

# Elections by proof of stake

## Basics

We assume a blockchain system in which  $N$  **secure tokens** are used:

- Each token has a unique **owner**
- Each token has a uniquely associated **index**  $1 \leq k \leq N$
- A token cannot be modified or copied without this going unnoticed

## Principle

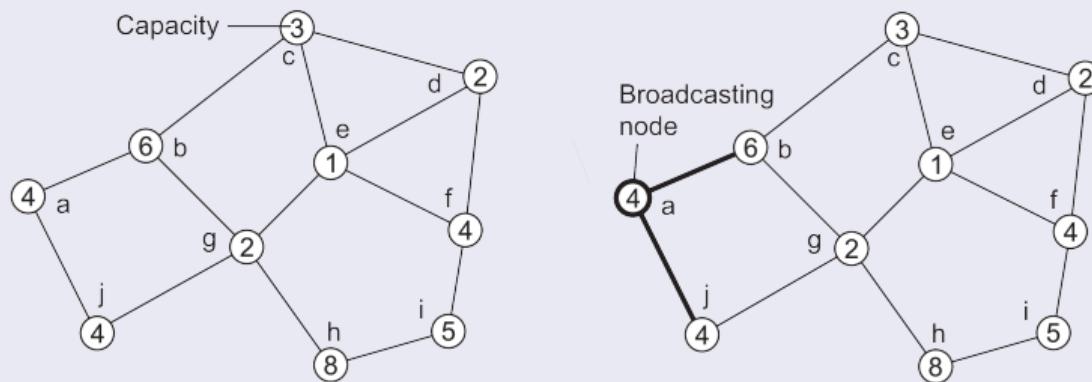
- Draw a random number  $k \in \{1, \dots, N\}$
- Look up the process  $P$  that owns the token with index  $k$ .  $P$  is the next leader.

## Observation

The more tokens a process owns, the higher the probability it will be selected as leader.

# A solution for wireless networks

## A sample network

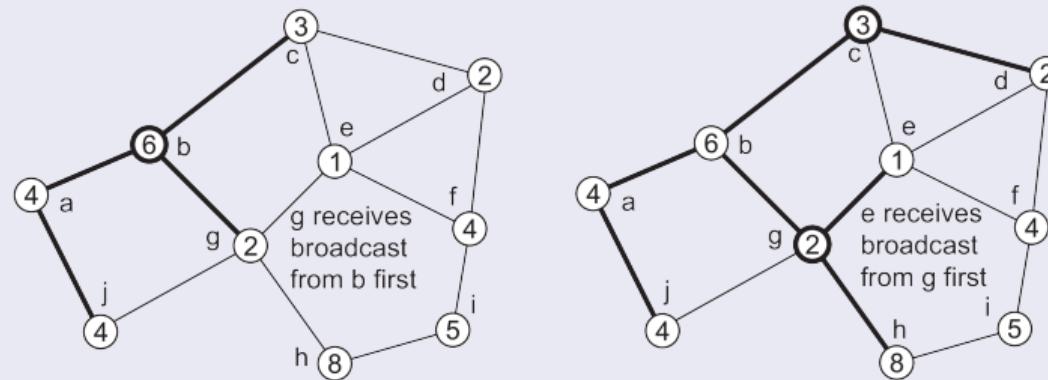


## Essence

Find the node with the highest capacity to select as the next leader.

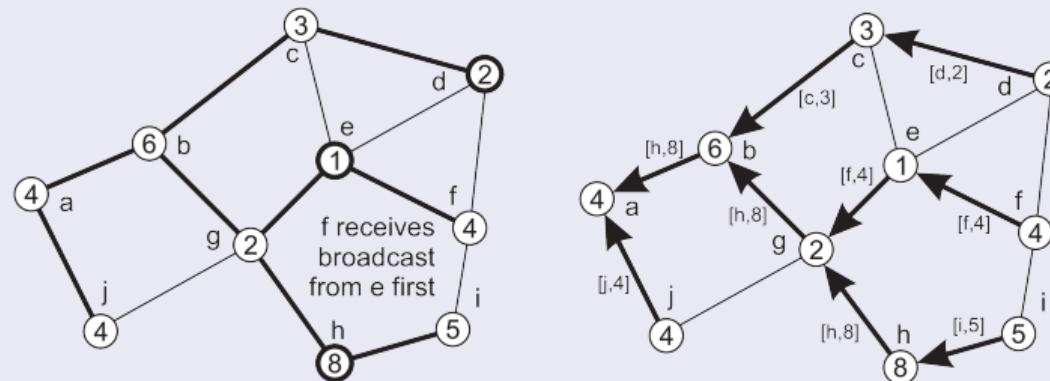
# A solution for wireless networks

## A sample network



# A solution for wireless networks

## A sample network



## Essence

A node reports back only the node that it found to have the highest capacity.

## Gossip-based coordination

## Gossip-based coordination

- Nodes randomly select a few other nodes (peers) periodically and exchange state information (e.g., membership, failure detection, key-value updates, or metadata).
- Information spreads exponentially fast, similar to how gossip or a virus spreads in a population, achieving eventual consistency with very high probability in logarithmic rounds.
- Advantages: Extremely robust to network partitions, node failures, and churn; no single point of failure; low coordination overhead; scales to thousands or tens of thousands of nodes with minimal latency impact; self-healing membership and failure detection (e.g., SWIM, HyParView, or Lifeguard protocols).
- Trade-offs & Common Uses: Provides only probabilistic guarantees (not strict consistency); latency of propagation is  $O(\log N)$ ; widely used for failure detection (accrual detectors),
- Lightweight broadcast instead of heavy consensus protocols like Paxos or Raft when strong consistency isn't required.

# Gossip-based coordination: aggregation

## Typical apps

- **Data dissemination**: Perhaps the most important one. Note that there are many variants of dissemination.
- **Aggregation Function (Average)**: Let every node  $P_i$  maintain a variable  $v_i$ . When two nodes gossip, they each reset their variable to

$$v_i, v_j \leftarrow (v_i + v_j)/2$$

Result: in the end each node will have computed the average  $\bar{v} = \sum_i v_i / N$ .

- What happens in the case that initially  $v_i = 1$  and  $v_j = 0, j \neq i$ ?

# Gossip-based coordination: peer sampling

## Problem

For many gossip-based applications, you need to **select a peer uniformly at random** from the entire network. In principle, this means you need to know all other peers. **Impossible?**

## Basics

- Each node maintains a list of  $c$  references to other nodes
- **Regularly**, pick another node at random (from the list), and **exchange** roughly  $c/2$  references
- When the **application** needs to select a node at random, it also picks a random one from its local list.

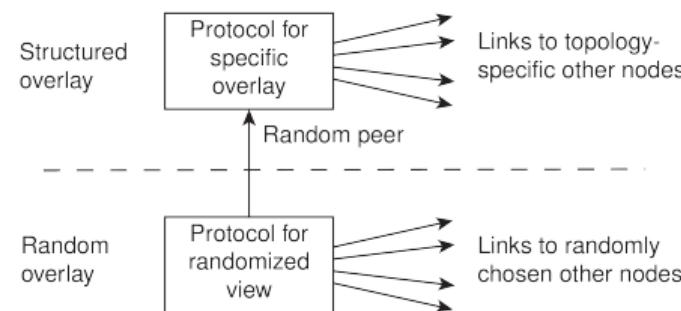
## Observation

Statistically, it turns out that the selection of a peer from the local list is indistinguishable from selecting uniformly at random peer from the entire network

# Gossip-based overlay construction

## Essence

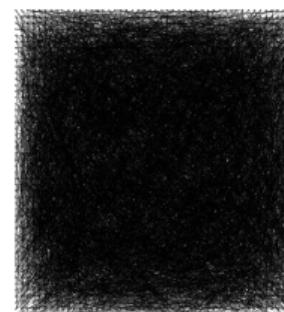
Maintain two local lists of neighbors. The lowest is used for providing a peer-sampling service; the highest list is used to carefully select application-dependent neighbors.



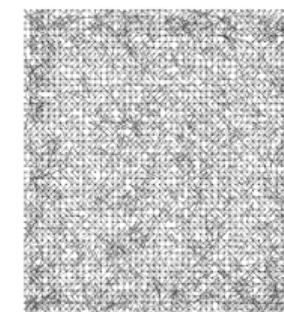
# Gossip-based overlay construction: a 2D torus

Consider a logical  $N \times N$  grid, with a node on each point of the grid.

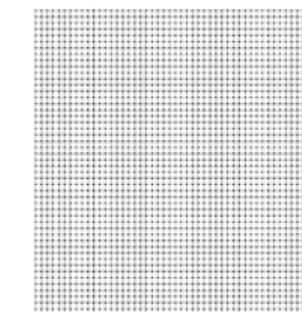
- Every node must maintain a list of  $c$  nearest neighbors
- Distance between node at  $(a_1, a_2)$  and  $(b_1, b_2)$  is  $d_1 + d_2$ , with  $d_i = \min(N - |a_i - b_i|, |a_i - b_i|)$
- Every node picks a random other node from its lowest-level list, and keeps only the closest one in its top-level list.
- Once every node has picked and selected a random node, we move to the next **round**



start ( $N = 50$ )



after 5 rounds

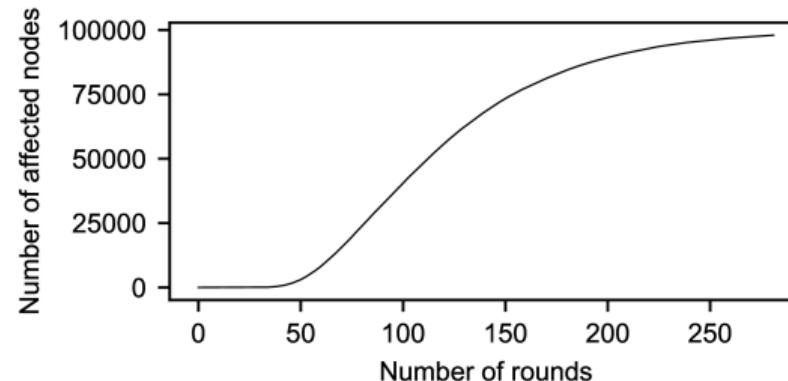


after 20 rounds

# Secure gossiping

## Dramatic attack

Consider when exchanging references, a set of [colluding nodes](#) systematically returns links only to each other  $\Rightarrow$  we are dealing with **hub attack**.

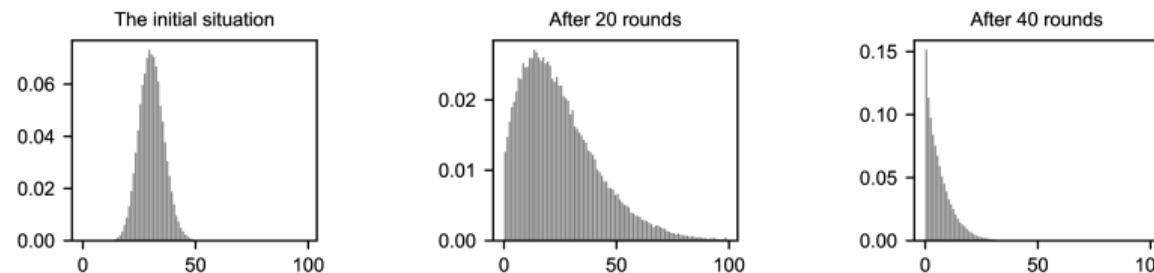


## Situation

A network with 100,000 nodes, a local list size  $c = 30$ , and only 30 attackers. The y-axis shows the number of nodes with links [only](#) to the attackers. After less than 300 rounds, the attackers have full control.

# A solution: gathering statistics

This is what measuring indegree distributions tells us: which fraction of nodes (y-axis) have how many other nodes pointing to them (x-axis)?



## Basic approach

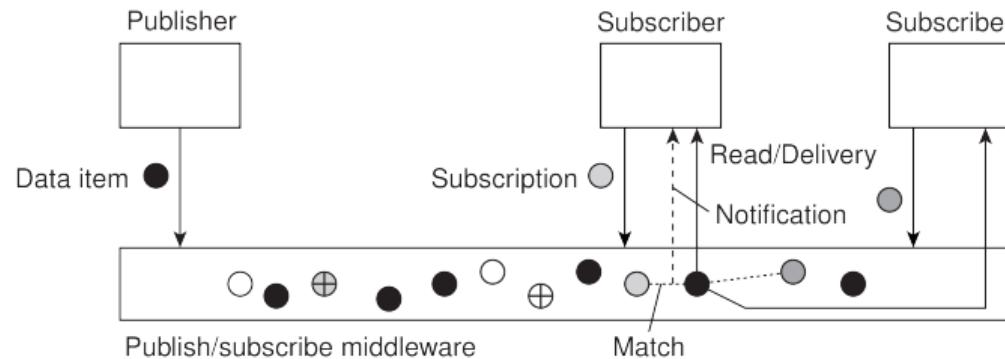
When a benign node initiates an exchange, it may either use the result for gathering statistics, or for updating its local list. An attacker is in limbo: will its response be used for statistical purposes or for functional purposes?

## Observation

When gathering statistics may reveal colluders, a colluding node will be **forced** to behave according to the protocol.

# Distributed event matching

# Distributed event matching



## Principle

- A process specifies in which events it is interested (subscription  $S$ )
- When a process publishes a notification  $N$  we need to see whether  $S$  matches  $N$ .

## Hard part

Implementing the **match** function in a scalable manner.

# General approach

## What is needed

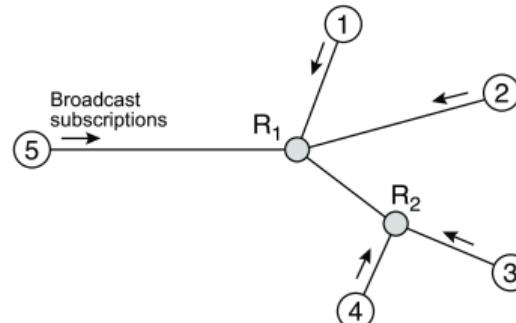
- $\text{sub2node}(S)$ : map a subscription  $S$  to a nonempty subset  $\mathbf{S}$  of servers
- $\text{not2node}(N)$ : map a notification  $N$  to a nonempty subset  $\mathbf{N}$  of servers

Make sure that  $\mathbf{S} \cap \mathbf{N} \neq \emptyset$ .

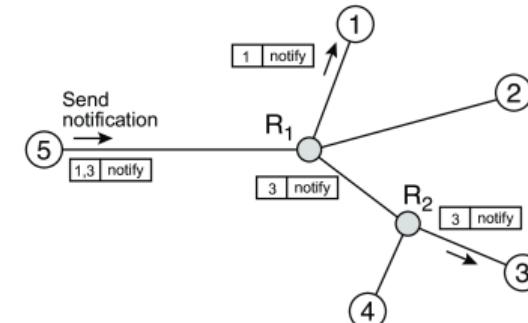
## Observations

- Centralized solution is simple:  $\mathbf{S} = \mathbf{N} = \{s\}$ , i.e. a single server.
- Topic-based publish-subscribe is also simple: each  $S$  and  $N$  is tagged with a **single topic**; each topic is handled by a single server (a **rendezvous node**). Several topics may be handled by same server).
- Content-based publish-subscribe is **tough**: a subscription takes the form **(attribute, value)** pair, with example values:
  - **range**: " $1 \leq x < 10$ "
  - **containment**: " $x \in \{\text{red}, \text{blue}\}$ "
  - **prefix and suffix expressions**: "`url.startswith("https")`"

# Selective routing



(a)



(b)

- (a) first broadcast subscriptions
- (b) forward notifications only to relevant rendezvous nodes

Example of a (partially filled) routing table

Interface	Filter
To node 3	$a \in [0, 3]$
To node 4	$a \in [2, 5]$
Toward router $R_1$	(unspecified)

# Gossiping: Sub-2-Sub

## Basics

- **Goal:** To realize scalability, make sure that subscribers with the same interests form just a single group
- **Model:** There are  $N$  attributes  $a_1, \dots, a_N$ . An attribute value is always (mappable to) a floating-point number.
- **Subscription:** Takes forms such as  $S = \langle a_1 \rightarrow 3.0, a_4 \rightarrow [0.0, 0.5] \rangle$ :  $a_1$  should be 3.0;  $a_4$  should lie between 0.0 and 0.5; other attribute values don't matter.

## Observations

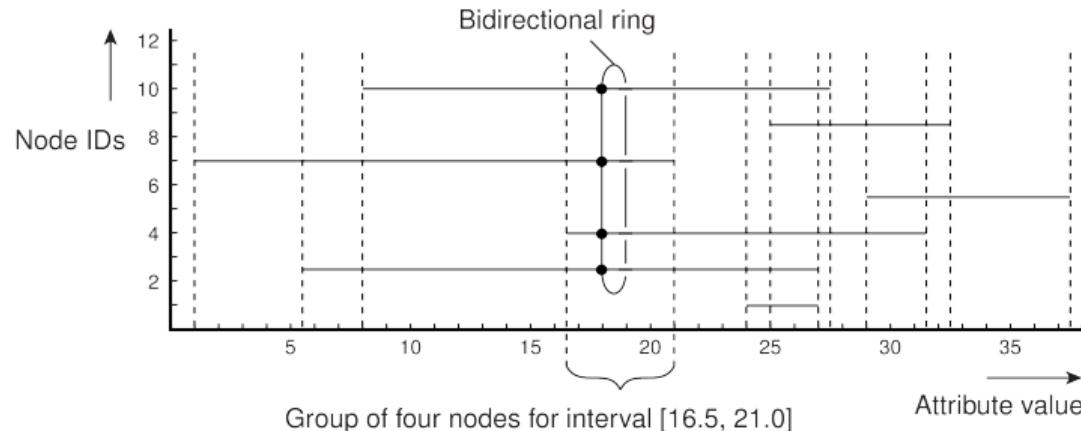
- A subscription  $S_i$  specifies a subset  $S_i$  in a  $N$ -dimensional space.
- We are interested only in notifications that fall into  $\bar{S} = \bigcup S_i$ .

## Goal

Partition  $\bar{S}$  into  $M$  disjoint subspaces  $\bar{S}_1, \dots, \bar{S}_M$  such that

- **Partitioning:**  $\forall k \neq m : \bar{S}_k \cap \bar{S}_m = \emptyset$  and  $\bigcup_m \bar{S}_m = \bar{S}$
- **Subscription coverage:**  $(\bar{S}_m \cap S_i \neq \emptyset) \Rightarrow (\bar{S}_m \subseteq S_i)$

# Gossiping: Sub-2-Sub



Consider a single attribute

- Nodes regularly exchange their subscriptions through gossiping
- An intersection between two nodes leads to a mutual reference
- If  $S_{ijk} = S_i \cap S_j \cap S_k \neq \emptyset$  and  $S_{ij} - S_{ijk} \neq \emptyset$ , then:
  - nodes  $i, j, k$  are grouped into a **single overlay network** (for  $S_{ijk}$ )
  - nodes  $i, j$  are grouped into a **single overlay network** (for  $S_{ij} - S_{ijk}$ )

# Secure publish-subscribe

We are facing nasty dilemma's

- **Referential decoupling:** messages should be able to flow from a publisher to subscribers while guaranteeing mutual anonymity ⇒ we cannot set up a secure channel.
- Not knowing where messages come from imposes **integrity problems**.
- Assuming a **trusted broker** may easily be practically impossible, certainly when dealing with sensitive information ⇒ we now have a **routing problem**.

This is a real concern when the broker / message bus is untrusted or semi-trusted (e.g., cloud provider, third-party MQTT broker, Kafka cluster you don't fully control)

# Secure publish-subscribe

Some real world examples:

- Financial market data feeds: Bloomberg, Refinitiv, or cloud feeds cannot let the provider see which banks are interested in “Argentina default”
- Privacy-preserving: ride-sharing / location pub-subDrivers publish encrypted location blocks; riders only want their 1 km grid square
- Military / classified pub-sub: Topics themselves are classified (**DRONE-STRIKE-X17**)

**We must use secure pub/sub whenever:**

- The broker is not fully trusted with message confidentiality, and
- Subscriptions are fine-grained (keywords, topics, attributes), and
- Publishers and subscribers are not in the same closed trust domain (open or Internet-facing system).

# Encrypted Everything?

Why "just encrypt everything" doesn't work If you encrypt messages (using asymmetric key encryption):

- Broker can't read them
- But the broker also can't route them, because:
- It doesn't know which subscription matches which message
- So the broker either:
- Sees everything (bad for security), or
- Sees nothing (bad for routing)

This is the core dilemma.

# Use Encrypted Search

## Solution

- Allow for searching (and matching) on encrypted data, without the need for decryption.
- **PEKS**: accompany encrypted messages with a collection of (again encrypted) keywords and search for matches on keywords.

Subscriber:

1. has private key  $sk$
2. decides he wants all documents containing "confidential"
3. computes  $T = \text{Trapdoor}(sk, \text{"confidential"})$
4. sends  $T$  (in the clear, over TLS) to the server/cloud

# PEKS

## Broker (Router/Server) tasks:

```
for each stored ciphertext C_i:  
    if Test(public_key, C_i, trapdoor_bytes) == True:  
        return this document to the Subscriber
```

- No encryption of the trapdoor is required,
- No shared secret between the Subscriber and the server is needed beyond normal transport-layer security (TLS).
- In this way, **PEKS** trapdoors are passed in all real systems.

# Summary

# Summary

The topics discussed in the *coordination* section of the lecture notes include

- ① Mutual Exclusion
- ② Election Algorithms
- ③ Gossip-based Coordination
- ④ Distributed Event Matching
- ⑤ Gossip-based Pub-Sub
- ⑥ Secure Gossip-based Pub-Sub