

# Principles of Distributed Systems

inft-3507

Dr. J.Burns

**ADA University**

Autumn 2025

## **Section 8: Fault Tolerance**

*This content is based on the following public resources: <https://www.distributed-systems.net/index.php/books/ds4/>*

# Introduction to fault tolerance

# Dependability

## Basics

A **component** provides **services** to **clients**. To provide services, the component may require the services from other components  $\Rightarrow$  a component may **depend** on some other component.

## Specifically

A component  $C$  depends on  $C^*$  if the **correctness** of  $C$ 's behavior depends on the correctness of  $C^*$ 's behavior. (Components are processes or channels.)

## Requirements related to dependability

Requirement	Description
<b>Availability</b>	Readiness for usage
<b>Reliability</b>	Continuity of service delivery
<b>Safety</b>	Very low probability of catastrophes
<b>Maintainability</b>	How easy can a failed system be repaired

# Reliability versus availability

## Reliability $R(t)$ of component $C$

Conditional probability that  $C$  has been functioning correctly during  $[0, t)$  given  $C$  was functioning correctly at time  $T = 0$ .

## Traditional metrics

- **Mean Time To Failure** ( $MTTF$ ): The average time until a component fails.
- **Mean Time To Repair** ( $MTTR$ ): The average time needed to repair a component.
- **Mean Time Between Failures** ( $MTBF$ ): Simply  $MTTF + MTTR$ .

# Reliability versus availability

## Availability $A(t)$ of component $C$

Average fraction of time that  $C$  has been up-and-running in interval  $[0, t)$ .

- Long-term availability  $A$ :  $A(\infty)$
- **Note:**  $A = \frac{MTTF}{MTBF} = \frac{MTTF}{MTTF + MTTR}$

## Observation

Reliability and availability make sense only if we have an accurate notion of what a failure actually is.

# Terminology

## Failure, error, fault

Term	Description	Example
Failure	A component is not living up to its specifications	Crashed program
Error	Part of a component that can lead to a failure	Programming bug
Fault	Cause of an error	Sloppy programmer

# Terminology

## Handling faults

Term	Description	Example
Fault prevention	Prevent the occurrence of a fault	Don't hire sloppy programmers
Fault tolerance	Build a component such that it can mask the occurrence of a fault	Build each component by two independent programmers
Fault removal	Reduce the presence, number, or seriousness of a fault	Get rid of sloppy programmers
Fault forecasting	Estimate current presence, future incidence, and consequences of faults	Estimate how a recruiter is doing when it comes to hiring sloppy programmers

# Failure models

## Types of failures

Type	Description of server's behavior
<b>Crash failure</b>	Halts, but is working correctly until it halts
<b>Omission failure</b> <i>Receive omission</i> <i>Send omission</i>	Fails to respond to incoming requests Fails to receive incoming messages Fails to send messages
<b>Timing failure</b>	Response lies outside a specified time interval
<b>Response failure</b> <i>Value failure</i> <i>State-transition failure</i>	Response is incorrect The value of the response is wrong Deviates from the correct flow of control
<b>Arbitrary failure</b>	May produce arbitrary responses at arbitrary times



# Dependability versus security

## Omission versus commission

Arbitrary failures are sometimes qualified as **malicious**. It is better to make the following distinction:

- **Omission failures**: a component fails to take an action that it should have taken
- **Commission failure**: a component takes an action that it should not have taken

## Observation

Note that **deliberate** failures, be they omission or commission failures, are typically security problems. Distinguishing between deliberate failures and unintentional ones is, in general, impossible.

# Halting failures

## Scenario

$C$  no longer perceives any activity from  $C^*$  — a **halting failure**? Distinguishing between a **crash** or **omission/timing failure** may be impossible.

## Asynchronous versus synchronous systems

- **Asynchronous system**: no assumptions about process execution speeds or message delivery times → cannot reliably detect crash failures.
- **Synchronous system**: process execution speeds and message delivery times are bounded → we can reliably detect omission and timing failures.
- In practice we have **partially synchronous systems**: most of the time, we can assume the system to be synchronous, yet there is no bound on the time that a system is asynchronous → can normally reliably detect crash failures.

# Halting failures

## Assumptions we can make

Halting type	Description
Fail-stop	Crash failures, but reliably detectable
Fail-noisy	Crash failures, eventually reliably detectable
Fail-silent	Omission or crash failures: clients cannot tell what went wrong
Fail-safe	Arbitrary, yet benign failures (i.e., they cannot do any harm)
Fail-arbitrary	Arbitrary, with malicious failures

# Redundancy for failure masking

## Types of redundancy


- **Information redundancy:** Add extra bits to data units so that errors can be recovered when bits are garbled.
- **Time redundancy:** Design a system such that an action can be performed again if anything went wrong. Typically used when faults are transient or intermittent.
- **Physical redundancy:** add equipment or processes in order to allow one or more components to fail. This type is extensively used in distributed systems.

## Process resilience

# Process resilience

## Basic idea

Protect against malfunctioning processes through **process replication**, organizing multiple processes into a **process group**. Distinguish between **flat groups** and **hierarchical groups**.



# Groups and failure masking

## $k$ -fault tolerant group

When a group can mask any  $k$  concurrent member failures ( $k$  is called **degree of fault tolerance**).

## How large does a $k$ -fault tolerant group need to be?

- With **halting failures** (crash/omission/timing failures): we need a total of  $k + 1$  members as **no member will produce an incorrect result, so the result of one member is good enough**.
- With **arbitrary failures**: we need  $2k + 1$  members so that the correct result can be obtained through a majority vote.

## Important assumptions

- All members are identical
- All members process commands in the same order

**Result:** We can now be sure that all processes do exactly the same thing.

# Consensus

## Prerequisite

In a fault-tolerant process group, each nonfaulty process executes the same commands, and in the same order, as every other nonfaulty process.

## Reformulation

Nonfaulty group members need to reach **consensus** on which command to execute next.



# Flooding-based consensus


## System model

- A process group  $\mathbf{P} = \{P_1, \dots, P_n\}$
- **Fail-stop** failure semantics, i.e., with **reliable failure detection**
- A client contacts a  $P_i$  requesting it to execute a command
- Every  $P_i$  maintains a list of proposed commands

## Basic algorithm (based on rounds)

- 1 In **round**  $r$ ,  $P_i$  multicasts its known set of commands  $\mathbf{C}_i^r$  to all others
- 2 At the end of  $r$ , each  $P_i$  merges all received commands into a new  $\mathbf{C}_i^{r+1}$ .
- 3 Next command  $cmd_i$  selected through a **globally shared, deterministic function**:  
 $cmd_i \leftarrow select(\mathbf{C}_i^{r+1})$ .

# Flooding-based consensus: Example



## Observations

- $P_2$  received all proposed commands from all other processes  $\Rightarrow$  **makes decision**.
- $P_3$  may have detected that  $P_1$  crashed, but does not know if  $P_2$  received anything, i.e.,  $P_3$  cannot know **if it has the same information** as  $P_2 \Rightarrow$  **cannot make decision** (same for  $P_4$ ).

# Raft

## Developed for understandability

- Uses a fairly straightforward **leader-election** algorithm (see Chp. 5). The current leader operates during the **current term**.
- Every server (typically, five) keeps a **log** of operations, some of which have been committed. **A backup will not vote for a new leader if its own log is more up to date.**
- All committed operations have the same position in the log of each respective server.
- The leader decides which pending operation is to be committed next  $\Rightarrow$  a **primary-backup approach**.

# Raft


## When submitting an operation

- A client submits a request for operation  $o$ .
- The leader appends the request  $\langle o, t, \rangle$  to its own log (registering the current term  $t$  and length of  $o$ ).
- The log is (conceptually) broadcast to the other servers.
- The others (conceptually) copy the log and acknowledge the receipt.
- When a majority of acks arrives, the leader commits  $o$ .

## Note

In practice, only updates are broadcast. At the end, every server has the same view and knows about the  $c$  committed operations. Note that effectively, any information at the backups is overwritten.

# Raft: when a leader crashes



## Crucial observations


- The new leader has the most committed operations in its log.
- Any missing commits will eventually be sent to the other backups.

## Consensus in faulty systems with arbitrary failures


# Consensus under arbitrary failure semantics

## Essence

We consider process groups in which communication between process is **inconsistent**.



Improper forwarding



Different messages

# Consensus under arbitrary failure semantics

## System model

- We consider a **primary**  $P$  and  $n - 1$  **backups**  $B_1, \dots, B_{n-1}$ .
- A client sends  $v \in \{T, F\}$  to  $P$
- Messages may be **lost**, but this can be detected.
- Messages **cannot be corrupted** beyond detection.
- A receiver of a message can **reliably detect its sender**.

## Byzantine agreement: requirements

**BA1:** Every nonfaulty backup process stores the same value.

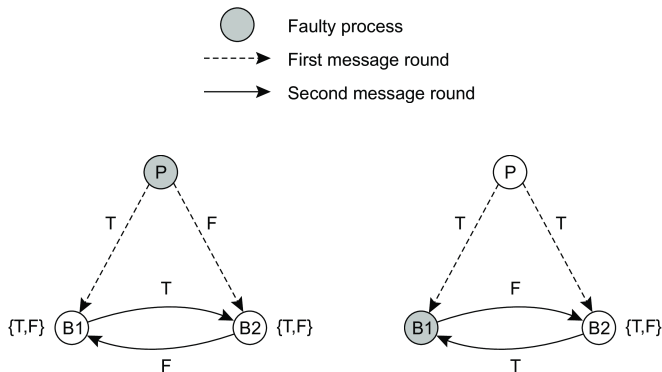
**BA2:** If the primary is nonfaulty then every nonfaulty backup process stores exactly what the primary had sent.

## Observation

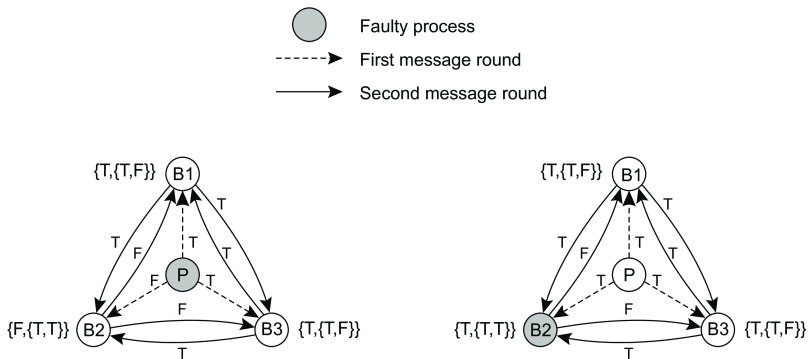
- Primary faulty  $\Rightarrow$  BA1 says that backups may store the same, but different (and thus wrong) value than originally sent by the client.
- Primary not faulty  $\Rightarrow$  satisfying BA2 implies that BA1 is satisfied.



# Why having $3k$ processes is not enough



# Why having $3k+1$ processes is enough



# Practical Byzantine Fault Tolerance (PBFT)

## Background

One of the first solutions that managed to Byzantine fault tolerance while keeping performance acceptable. Popularity has increased with the introduction of [permissioned blockchains](#).


## Assumptions

- A server may exhibit arbitrary failures
- Messages may be lost, delayed, and received out of order
- Messages have an [identifiable sender](#) (i.e., they are [signed](#))
- [Partially synchronous](#) execution model

## Essence


A [primary-backup approach](#) with  $3k + 1$  replica servers.

# PBFT: four phases




- C is the client
- P is the primary
- B<sub>1</sub>, B<sub>2</sub>, B<sub>3</sub> are backups
- Assume B<sub>2</sub> is faulty

# PBFT: four phases




- All servers assume to be working in a current **view**  $v$ .
- C requests operation  $o$  to be executed
- P **timestamps**  $o$  and sends **PRE-PREPARE**( $t, v, o$ )
- Backup  $B_i$  accepts the pre-prepare message if it is also in  $v$  and has not accepted a an operation with timestamp  $t$  before.

# PBFT: four phases



- $B_i$  broadcasts  $\text{PREPARE}(t, v, o)$  to all (including the primary)
- **Note:** a nonfaulty server will eventually log  $2k$  messages  $\text{PREPARE}(t, v, o)$  (including its own)  $\Rightarrow$  consensus on the ordering of  $o$ .
- **Note:** it doesn't matter what faulty  $B_2$  sends, it cannot affect joint decisions by  $P$ ,  $B_1$ ,  $B_3$ .

# PBFT: four phases



- All servers broadcast **COMMIT**( $t, v, o$ )
- The commit is needed to also make sure that  $o$  can be executed **now**, that is, in the current view  $v$ .
- When  $2k$  messages have been collected, excluding its own, the server can safely execute  $o$  en reply to the client.

# PBFT: when the primary fails

## Issue

When a backup detects the primary failed, it will broadcast a **view change** to view  $v+1$ . We need to ensure that any **outstanding request** is executed **once and only once** by all nonfaulty servers. The operation needs to be handed over to the new view.

## Procedure

- The next primary  $P^*$  is known deterministically
- A backup server broadcasts **VIEW-CHANGE**( $v+1, \mathbf{P}$ ):  $\mathbf{P}$  is the set of prepares it had sent out.
- $P^*$  waits for  $2k+1$  view-change messages, with  $\mathbf{X} = \bigcup \mathbf{P}$  containing all previously sent prepares.
- $P^*$  sends out **NEW-VIEW**( $v+1, \mathbf{X}, \mathbf{O}$ ) with  $\mathbf{O}$  a new set of pre-prepare messages.
- **Essence**: this allows the nonfaulty backups to **replay** what has gone on in the previous view, if necessary, and bring  $o$  into the new view  $v+1$ .



# Realizing fault tolerance

## Observation

Considering that the members in a fault-tolerant process group are so tightly coupled, we may bump into considerable performance problems, but perhaps even situations in which realizing fault tolerance is impossible.

## Question

Are there limitations to what can be readily achieved?

- What is needed to enable reaching consensus?
- What happens when groups are partitioned?

# Distributed consensus: when can it be reached

Process behavior	Message ordering				Commun. delay
	Unordered		Ordered		
	Unicast	Multicast	Unicast	Multicast	
Synchronous	✓	✓	✓	✓	Bounded
			✓	✓	Unbounded
Asynchronous				✓	Bounded
				✓	UnBounded
	Unicast	Multicast	Unicast	Multicast	
Message transmission					

## Formal requirements for consensus

- Processes produce the same output value
- Every output value must be valid
- Every process must eventually provide output

# Consistency, availability, and partitioning

## CAP theorem

Any networked system providing shared data can provide only two of the following three properties:

- C:** **consistency**, by which a shared and replicated data item appears as a single, up-to-date copy
- A:** **availability**, by which updates will always be eventually executed
- P:** Tolerant to the **partitioning** of process group.

## Conclusion

In a network subject to communication failures, it is impossible to realize an atomic read/write **shared memory** that guarantees a response to every request.

# CAP theorem intuition

## Simple situation: two interacting processes

- $P$  and  $Q$  can no longer communicate:
  - Allow  $P$  and  $Q$  to go ahead  $\Rightarrow$  no consistency
  - Allow only one of  $P$ ,  $Q$  to go ahead  $\Rightarrow$  no availability
- $P$  and  $Q$  have to be assumed to continue communication  $\Rightarrow$  no partitioning allowed.

## Fundamental question

What are the practical ramifications of the CAP theorem?

# Failure detection

## Issue

How can we **reliably detect** that a process has **actually crashed**?

## General model

- Each process is equipped with a failure detection module
- A process  $P$  **probes** another process  $Q$  for a reaction
- If  $Q$  reacts:  $Q$  is considered to be alive (by  $P$ )
- If  $Q$  does not react with  $t$  time units:  $Q$  is **suspected** to have crashed

## Observation for a **synchronous** system

a suspected crash  $\equiv$  a known crash

# Practical failure detection

## Implementation

- If  $P$  did not receive **heartbeat** from  $Q$  within time  $t$ :  $P$  **suspects**  $Q$ .
- If  $Q$  later sends a message (which is received by  $P$ ):
  - $P$  stops suspecting  $Q$
  - $P$  increases the timeout value  $t$
- **Note:** if  $Q$  did crash,  $P$  will keep suspecting  $Q$ .

# Recovery

# Recovery: Background

## Essence

When a failure occurs, we need to bring the system into an error-free state:

- **Forward error recovery:** Find a new state from which the system can continue operation
- **Backward error recovery:** Bring the system back into a **previous** error-free state

## Practice

Use backward error recovery, requiring that we establish **recovery points**

## Observation

Recovery in distributed systems is complicated by the fact that processes need to cooperate in identifying a **consistent state** from where to recover




# Consistent recovery state

## Requirement

Every message that has been received is also shown to have been sent in the state of the sender.

## Recovery line

Assuming processes regularly **checkpoint** their state, the most recent **consistent global checkpoint**.



# Coordinated checkpointing

## Essence

Each process takes a checkpoint after a globally coordinated action.

## Simple solution

Use a two-phase blocking protocol:

- A coordinator multicasts a **checkpoint request** message
- When a participant receives such a message, it takes a checkpoint, stops sending (application) messages, and reports back that it has taken a checkpoint
- When all checkpoints have been confirmed at the coordinator, the latter broadcasts a **checkpoint done** message to allow all processes to continue


## Observation

It is possible to consider only those processes that depend on the recovery of the coordinator, and ignore the rest

# Cascaded rollback

## Observation

If checkpointing is done at the “wrong” instants, the recovery line may lie at system startup time. We have a so-called **cascaded rollback**.



# Independent checkpointing

## Essence

Each process independently takes checkpoints, with the risk of a cascaded rollback to system startup.

- Let  $CP_i(m)$  denote  $m^{\text{th}}$  checkpoint of process  $P_i$  and  $INT_i(m)$  the interval between  $CP_i(m-1)$  and  $CP_i(m)$ .
- When process  $P_i$  sends a message in interval  $INT_i(m)$ , it piggybacks  $(i, m)$
- When process  $P_j$  receives a message in interval  $INT_j(n)$ , it records the dependency  $INT_i(m) \rightarrow INT_j(n)$ .
- The dependency  $INT_i(m) \rightarrow INT_j(n)$  is saved to storage when taking checkpoint  $CP_j(n)$ .

## Observation

If process  $P_i$  rolls back to  $CP_i(m-1)$ ,  $P_j$  must roll back to  $CP_j(n-1)$ .

# Message logging

## Alternative

Instead of taking an (expensive) checkpoint, try to **replay** your (communication) behavior from the most recent checkpoint  $\Rightarrow$  store messages in a log.

## Assumption

We assume a **piecewise deterministic** execution model:

- The execution of each process can be considered as a sequence of state intervals
- Each state interval starts with a nondeterministic event (e.g., message receipt)
- Execution in a state interval is deterministic

## Conclusion


If we record nondeterministic events (to replay them later), we obtain a deterministic execution model that will allow us to do a complete replay.

# Message logging and consistency

## When should we actually log messages?

Avoid **orphan processes**:

- Process  $Q$  has just received and delivered messages  $m_1$  and  $m_2$
- Assume that  $m_2$  is never logged.
- After delivering  $m_1$  and  $m_2$ ,  $Q$  sends message  $m_3$  to process  $R$
- Process  $R$  receives and subsequently delivers  $m_3$ : it is an orphan.



# Message-logging schemes

## Notations

- **DEP**( $m$ ): processes to which  $m$  has been delivered. If message  $m^*$  is causally dependent on the delivery of  $m$ , and  $m^*$  has been delivered to  $Q$ , then  $Q \in \mathbf{DEP}(m)$ .
- **COPY**( $m$ ): processes that have a copy of  $m$ , but have not (yet) reliably stored it.
- **FAIL**: the collection of crashed processes.

## Characterization

$Q$  is orphaned  $\Leftrightarrow \exists m : Q \in \mathbf{DEP}(m)$  and  $\mathbf{COPY}(m) \subseteq \mathbf{FAIL}$

# Message-logging schemes

## Pessimistic protocol

For each **nonstable** message  $m$ , there is at most one process dependent on  $m$ , that is  $|\mathbf{DEP}(m)| \leq 1$ .

## Consequence

An unstable message in a pessimistic protocol **must** be made stable before sending a next message.



# Message-logging schemes

## Optimistic protocol

For each unstable message  $m$ , we ensure that if **COPY**( $m$ )  $\subseteq$  **FAIL**, then eventually also **DEP**( $m$ )  $\subseteq$  **FAIL**.

## Consequence

To guarantee that **DEP**( $m$ )  $\subseteq$  **FAIL**, we generally roll back each orphan process  $Q$  until  $Q \notin \mathbf{DEP}(m)$ .

## Summary

# Summary

In this section on *Fault Tolerance*, we discussed the following key concepts:

- Dependability, Reliability, Availability
- Process Resilience
- Consensus in faulty systems with arbitrary failures
- Practical Byzantine Fault Tolerance (PBFT)
- Recovery