

# Principles of Distributed Systems

inft-3507

Dr. J.Burns

**ADA University**

Autumn 2025

## Section 3: Processes and Threads

*This content is based on the following public resources: <https://www.distributed-systems.net/index.php/books/ds4/>*

# Processes and Threads

# Introduction to Processes and Threads

- **Definition:** A process is an independent program in execution with its own memory space, while a thread is a lightweight unit of execution within a process, sharing the process's memory.
- **Memory:** A process has its own isolated memory space (address space), whereas threads within the same process share the same memory space, including code, data, and resources.
- **Overhead:** Processes are heavier, requiring more system resources and time for creation and context switching, while threads are lighter, with lower overhead for creation and switching.
- **Communication:** Inter-process communication (IPC) is complex and slower (e.g., pipes, sockets), while threads communicate faster via shared memory but require synchronization (e.g., locks).

# Context switching

## Observations

- ① Threads share the same address space. Thread context switching is much faster than process context switching:
  - ① only registers and program counter need to be saved and restored
- ② Process context switching is more expensive (in time and space) as
  - ① TLB needs to be flushed
  - ② page table is reloaded
  - ③ address space changes
- ③ Creating and destroying threads is much cheaper than doing so for processes.


# Why use threads

## Some simple reasons

- **Avoid needless blocking:** a single-threaded process will **block** when doing I/O; in a multithreaded process, the operating system can switch the CPU to another thread in that process.
- **Exploit parallelism:** the threads in a multithreaded process can be scheduled to run in parallel on a multiprocessor or multicore processor.
- **Avoid process switching:** structure large applications not as a collection of processes, but through multiple threads.

# Avoid process switching

## Avoid expensive context switching



## Trade-offs

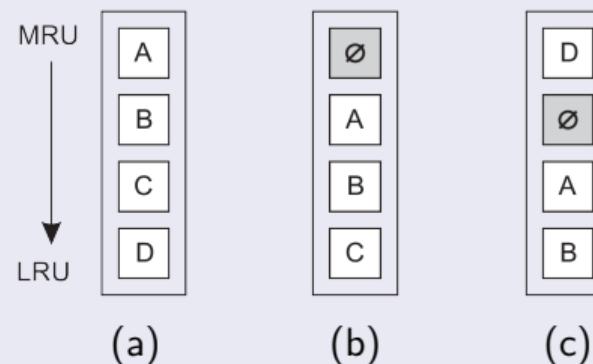
- Threads use the same address space: more prone to errors
- No support from OS/HW to protect threads using each other's memory
- Thread context switching may be faster than process context switching

# The cost of a context switch

Consider a simple clock-interrupt handler

- **direct costs**: actual switch and executing code of the handler
- **indirect costs**: other costs, notably caused by messing up the cache

What a context switch may cause: indirect costs



- (a) before the context switch  
(b) after the context switch  
(c) after accessing block *D*.

# Thread Switch vs Process Context Switch

## Thread Switch

- Switch between threads of the *same process*
- Address space remains unchanged
- Page tables and memory mappings are reused
- Faster than process switches
- Typically used for concurrency within applications

### Overhead:

- Save/restore registers
- Update stack pointer and program counter

## Process Context Switch

- Switch between threads of *different processes*
- Address space changes
- Page tables must be switched
- TLB often flushed or invalidated
- Slower than thread switches

### Overhead:

- Save/restore registers
- Switch address space
- Update memory management state

# Threads and operating systems

## Main issue

Should an OS kernel provide threads, or should they be implemented as user-level packages?

## User-space solution

- All operations can be completely handled **within a single process** ⇒ implementations can be extremely efficient.
- **All** services provided by the kernel are done **on behalf of the process in which a thread resides** ⇒ if the kernel decides to block a thread, the entire process will be blocked.
- Threads are used when there are many external events: **threads block on a per-event basis** ⇒ if the kernel can't distinguish threads, how can it support signaling events to them?

# Linux Kernel Threads

- **Task Struct Representation:** In the Linux kernel, threads are implemented as lightweight processes, each represented by a `task_struct` (defined in `include/linux/sched.h`), sharing memory but maintaining separate execution contexts for scheduling.
- **Scheduling with CFS:** The Completely Fair Scheduler (CFS) in `kernel/sched/fair.c` manages kernel threads, treating them as virtual processors and allocating CPU time fairly using a red-black tree.
- **POSIX Threads Integration:** User-space threads (e.g., via `pthread_create`) map to kernel threads, enabling Java's 1:1 threading model to leverage Linux's scheduling for efficient concurrency.

# Using threads at the client side

## Multithreaded web client

### Hiding network latencies:

- Web browser scans an incoming HTML page, and finds that **more files need to be fetched**.
- **Each file is fetched by a separate thread**, each doing a (blocking) HTTP request.
- As files come in, the browser displays them.

## Multiple request-response calls to other machines (RPC)

- A client does several calls at the same time, each one by a different thread.
- It then waits until all results have been returned.
- Note: if calls are to different servers, we may have a **linear speed-up**.

# Multithreaded clients: does it help?

## Thread-level parallelism: TLP

Let  $c_i$  denote the fraction of time that exactly  $i$  threads are being executed simultaneously.

$$TLP = \frac{\sum_{i=1}^N i \cdot c_i}{1 - c_0}$$

with  $N$  the maximum number of threads that (can) execute at the same time.

## Practical measurements

A typical Web browser has a TLP value between 1.5 and 2.5  $\Rightarrow$  threads are primarily used for logically organizing browsers.

# Using threads at the server side

## Improve performance


- Starting a thread is cheaper than starting a new process.
- Having a single-threaded server prohibits simple scale-up to a multiprocessor system.
- As with clients: [hide network latency](#) by reacting to next request while previous one is being replied.

## Better structure

- Most servers have high I/O demands. Using simple, [well-understood blocking calls](#) simplifies the structure.
- Multithreaded programs tend to be smaller and easier to understand due to simplified flow of control.

# Why multithreading is popular: organization

## Dispatcher/worker model



## Overview

Multithreading	Parallelism, blocking system calls
Single-threaded process	No parallelism, blocking system calls

# Virtualization

# Problem Statement and Scope

## Objective

Virtualization provides multiple isolated execution environments on shared hardware while preserving operating system semantics and acceptable performance.

- Motivation: consolidation, isolation, manageability, and cloud-scale deployment
- Focus: virtual machines, virtual machine monitors (VMMs), and hypervisors
- Perspective: contemporary hardware-assisted virtualization


# The Virtual Machine Abstraction

## Definition

A virtual machine is a software-defined abstraction of a physical computer, including CPU, memory, devices, and firmware.


- Presents a conventional hardware interface to the guest OS
- Preserves OS assumptions about privilege, memory, interrupts, and devices
- Enables independent failure, security, and resource control domains

# System Structure and Trusted Computing Base




**Figure:** Logical layering of a virtual machine system. The trusted computing base consists primarily of the hypervisor core and device model.

# CPU Virtualization: VM Entry and Exit




**Figure:** Hardware-supported transitions between guest and hypervisor execution contexts.

# Memory Virtualization Using Nested Paging



**Figure:** Two-dimensional address translation used in hardware-assisted memory virtualization.

# I/O Virtualization Design Space

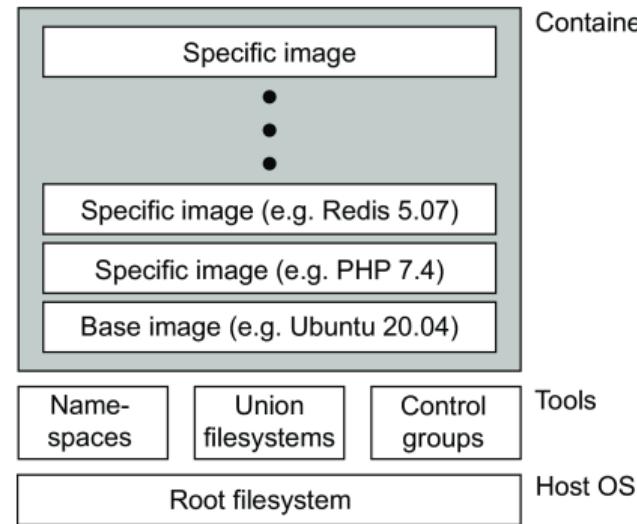


**Figure:** Spectrum of I/O virtualization techniques and their trade-offs.

# Current Directions in Virtualization

- MicroVMs reduce device models and trusted code size
- Confidential VMs protect memory from privileged software
- Virtual machines serve as isolation boundaries for containerized workloads
- Integration with cluster schedulers and cloud control planes

# Containers



- **Namespaces**: a collection of processes in a container is given their own view of identifiers
- **Union file system**: combine several file systems into a layered fashion with only the highest layer allowing for write operations (and the one being part of a container).
- **Control groups**: resource restrictions can be imposed upon a collection of processes.

# VMs and cloud computing

## Three types of cloud services

- Infrastructure-as-a-Service covering the basic infrastructure
- Platform-as-a-Service covering system-level services
- Software-as-a-Service containing actual applications

## IaaS

Instead of renting out a physical machine, a cloud provider will rent out a VM (or VMM) that may be sharing a physical machine with other customers ⇒ almost complete isolation between customers (although performance isolation may not be reached).

# Summary

# Summary and Conclusions

We have discussed processes and threads in Distributed Systems, namely:

- Processes and Threads
- Context Switching
- Multithreading
- Virtualization
- Containerization