

# Principles of Distributed Systems

inft-3507

Dr. J.Burns

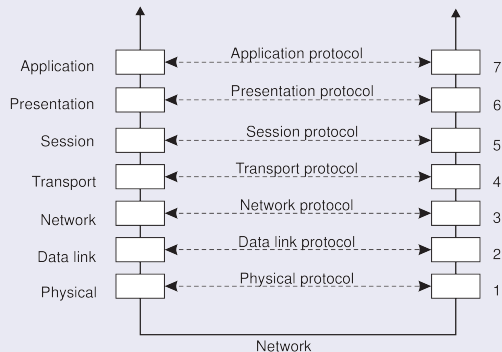
**ADA University**

Autumn 2025

## **Section 4: Communication**

Ref: <https://www.distributed-systems.net/index.php/books/ds4/>

# Basic networking model



## Drawbacks

- Focus on message-passing only
- Often unneeded or unwanted functionality
- Violates access transparency

# Low-level layers

## Recap

- **Physical layer**: contains the specification and implementation of bits, and their transmission between sender and receiver
- **Data link layer**: prescribes the transmission of a series of bits into a frame to allow for error and flow control
- **Network layer**: describes how packets in a network of computers are to be **routed**.

## Observation

For many distributed systems, the lowest-level interface is that of the network layer.

# Transport Layer

## Important

The transport layer provides the actual communication facilities for most distributed systems.

## Standard Internet protocols

- TCP: connection-oriented, reliable, stream-oriented communication
- UDP: unreliable (best-effort) datagram communication

# Middleware layer

## Observation


Middleware is invented to provide **common** services and protocols that can be used by many **different** applications

- A rich set of **communication protocols**
- **(Un)marshaling** of data, necessary for integrated systems
- **Naming protocols**, to allow easy sharing of resources
- **Security protocols** for secure communication
- **Scaling mechanisms**, such as for replication and caching

## Note

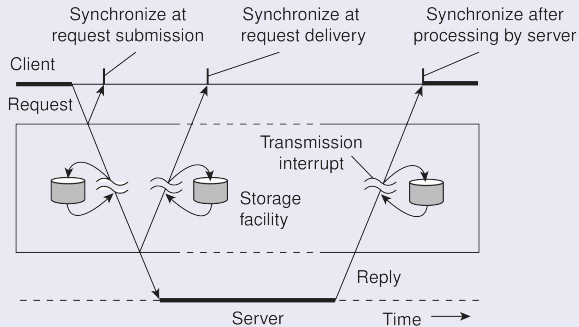
What remains are truly **application-specific** protocols... **such as?**

# An adapted layering scheme



# Types of communication


## Distinguish...



- **Transient** versus **persistent** communication
- **Asynchronous** versus **synchronous** communication

# Types of communication

## Transient versus persistent




- **Transient communication:** Comm. server discards message when it cannot be delivered at the next server, or at the receiver.
- **Persistent communication:** A message is stored at a communication server as long as it takes to deliver it.



# Types of communication

## Places for synchronization



- At request submission
- At request delivery
- After request processing

# Client/Server

## Some observations

Client/Server computing is generally based on a model of **transient synchronous communication**:

- Client and server have to be active at the time of communication
- Client issues request and blocks until it receives reply
- Server essentially waits only for incoming requests, and subsequently processes them

# Client/Server

## Some observations

Client/Server computing is generally based on a model of **transient synchronous communication**:

- Client and server have to be active at the time of communication
- Client issues request and blocks until it receives reply
- Server essentially waits only for incoming requests, and subsequently processes them

## Drawbacks synchronous communication

- Client cannot do any other work while waiting for reply
- Failures have to be handled immediately: the client is waiting
- The model may simply not be appropriate (mail, news)

# Messaging

## Message-oriented middleware

Aims at high-level **persistent asynchronous communication**:

- Processes send each other messages, which are queued
- Sender need not wait for immediate reply, but can do other things
- Middleware often ensures fault tolerance


# Basic RPC operation

## Observations

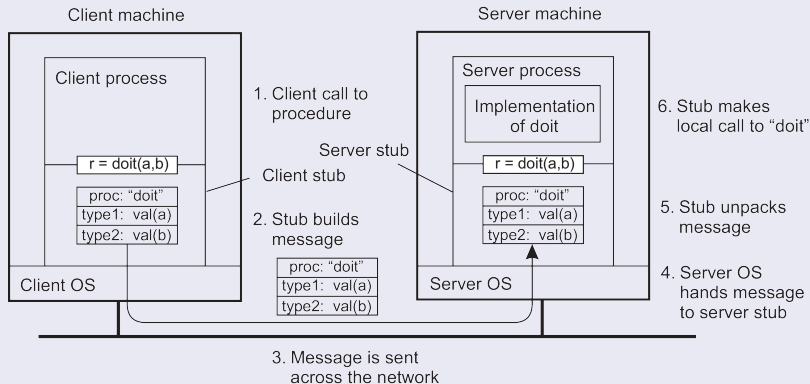
- Application developers are familiar with simple procedure model
- Well-engineered procedures operate in isolation (black box)
- There is no fundamental reason not to execute procedures on separate machine

## Conclusion

Communication between caller & callee can be hidden by using procedure-call mechanism.



# Basic RPC operation



- 1 Client procedure calls client stub.
- 2 Stub builds message; calls local OS.
- 3 OS sends message to remote OS.
- 4 Remote OS gives message to stub.
- 5 Stub unpacks parameters; calls server.

- 6 Server does local call; returns result to stub.
- 7 Stub builds message; calls OS.
- 8 OS sends message to client's OS.
- 9 Client's OS gives message to stub.
- 10 Client stub unpacks result; returns to client.

# RPC: Parameter passing

## There's more than just wrapping parameters into a message

- Client and server machines may have **different data representations** (think of byte ordering)
- Wrapping a parameter means **transforming a value into a sequence of bytes**
- Client and server have to **agree on the same encoding**:
  - How are **basic data values** represented (integers, floats, characters)
  - How are **complex data values** represented (arrays, unions)

## Conclusion

Client and server need to **properly interpret messages**, transforming them into machine-dependent representations.

# RPC: Parameter passing

## Some assumptions

- **Copy in/copy out** semantics: while procedure is executed, nothing can be assumed about parameter values.
- **All** data that is to be operated on is passed by parameters. Excludes passing **references to (global) data**.



# RPC: Parameter passing

## Some assumptions

- Copy in/copy out semantics: while procedure is executed, nothing can be assumed about parameter values.
- All data that is to be operated on is passed by parameters. Excludes passing references to (global) data.

## Conclusion

Full access transparency cannot be realized.

# RPC: Parameter passing

## Some assumptions

- **Copy in/copy out** semantics: while procedure is executed, nothing can be assumed about parameter values.
- **All** data that is to be operated on is passed by parameters. Excludes passing **references to (global) data**.

## Conclusion

Full access transparency cannot be realized.


## A **remote reference** mechanism enhances access transparency

- Remote reference offers **unified access** to remote data
- Remote references can be **passed as parameter** in RPCs
- **Note:** stubs can sometimes be used as such references

# Asynchronous RPCs

## Essence


Try to get rid of the strict request-reply behavior, but let the client continue without waiting for an answer from the server.



# Sending out multiple RPCs

## Essence


Sending an RPC request to a group of servers.



# Transient messaging: sockets

## Berkeley socket interface

Operation	Description
socket	Create a new communication end point
bind	Attach a local address to a socket
listen	Tell operating system what the maximum number of pending connection requests should be
accept	Block caller until a connection request arrives
connect	Actively attempt to establish a connection
send	Send some data over the connection
receive	Receive some data over the connection
close	Release the connection



# Sockets: Python code

## Server

```
1  from socket import *
2
3  class Server:
4      def run(self):
5          s = socket(AF_INET, SOCK_STREAM)
6          s.bind((HOST, PORT))
7          s.listen(1)
8          (conn, addr) = s.accept() # returns new socket and addr. client
9          while True: # forever
10             data = conn.recv(1024) # receive data from client
11             if not data: break # stop if client stopped
12             conn.send(data+b"*\") # return sent data plus an "*"
13             conn.close() # close the connection
```

## Client

```
1  class Client:
2      def run(self):
3          s = socket(AF_INET, SOCK_STREAM)
4          s.connect((HOST, PORT)) # connect to server (block until accepted)
5          s.send(b"Hello, world") # send same data
6          data = s.recv(1024) # receive the response
7          print(data) # print what you received
8          s.send(b"") # tell the server to close
9          s.close() # close the connection
```

# Making sockets easier to work with

## Observation

Sockets are rather low level and programming mistakes are easily made. However, the way that they are used is often the same (such as in a client-server setting).

## Alternative: ZeroMQ

Provides a higher level of expression by **pairing** sockets: one for sending messages at process  $P$  and a corresponding one at process  $Q$  for receiving messages. All communication is **asynchronous**.

## Three patterns

- Request-reply
- Publish-subscribe
- Pipeline

# Request-reply

```
1 import zmq
2
3 def server():
4     context = zmq.Context()
5     socket = context.socket(zmq.REP)      # create reply socket
6     socket.bind("tcp://*:12345")         # bind socket to address
7
8     while True:
9         message = socket.recv()           # wait for incoming message
10        if not "STOP" in str(message):    # if not to stop...
11            reply = str(message.decode())+'*' # append "*" to message
12            socket.send(reply.encode())     # send it away (encoded)
13        else:
14            break                          # break out of loop and end
15
16 def client():
17     context = zmq.Context()
18     socket = context.socket(zmq.REQ)      # create request socket
19
20     socket.connect("tcp://localhost:12345") # block until connected
21     socket.send(b"Hello world")           # send message
22     message = socket.recv()               # block until response
23     socket.send(b"STOP")                  # tell server to stop
24     print(message.decode())               # print result
```



# Publish-subscribe

```
1 import multiprocessing
2 import zmq, time
3
4 def server():
5     context = zmq.Context()
6     socket = context.socket(zmq.PUB)           # create a publisher socket
7     socket.bind("tcp://*:12345")              # bind socket to the address
8     while True:
9         time.sleep(5)                          # wait every 5 seconds
10        t = "TIME " + time.asctime()
11        socket.send(t.encode())                 # publish the current time
12
13 def client():
14     context = zmq.Context()
15     socket = context.socket(zmq.SUB)           # create a subscriber socket
16     socket.connect("tcp://localhost:12345")    # connect to the server
17     socket.setsockopt(zmq.SUBSCRIBE, b"TIME") # subscribe to TIME messages
18
19     for i in range(5):                         # Five iterations
20         time = socket.recv()                   # receive a message related to subscription
21         print(time.decode())                   # print the result
```

# Pipeline

```
1 def producer():
2     context = zmq.Context()
3     socket = context.socket(zmq.PUSH)      # create a push socket
4     socket.bind("tcp://127.0.0.1:12345")    # bind socket to address
5
6     while True:
7         workload = random.randint(1, 100)  # compute workload
8         socket.send(pickle.dumps(workload)) # send workload to worker
9         time.sleep(workload/NWORKERS)      # balance production by waiting
10
11 def worker(id):
12     context = zmq.Context()
13     socket = context.socket(zmq.PULL)      # create a pull socket
14     socket.connect("tcp://localhost:12345") # connect to the producer
15
16     while True:
17         work = pickle.loads(socket.recv())  # receive work from a source
18         time.sleep(work)                   # pretend to work
```

# MPI: When lots of flexibility is needed

## Representative operations

Operation	Description
MPI_BSEND	Append outgoing message to a local send buffer
MPI_SEND	Send a message and wait until copied to local or remote buffer
MPI_SSEND	Send a message and wait until transmission starts
MPI_SENDRECV	Send a message and wait for reply
MPI_ISEND	Pass reference to outgoing message, and continue
MPI_ISSEND	Pass reference to outgoing message, and wait until receipt starts
MPI_RECV	Receive a message; block if there is none
MPI_IRECV	Check if there is an incoming message, but do not block

# Queue-based messaging

## Four possible combinations

Sender  
running



Receiver  
running

Sender  
running



Receiver  
passive

Sender  
passive



Receiver  
running

Sender  
passive



Receiver  
passive

# Message-oriented middleware

## Essence

Asynchronous persistent communication through support of middleware-level **queues**. Queues correspond to buffers at communication servers.

## Operations


Operation	Description
PUT	Append a message to a specified queue
GET	Block until the specified queue is nonempty, and remove the first message
POLL	Check a specified queue for messages, and remove the first. Never block
NOTIFY	Install a handler to be called when a message is put into the specified queue

# General model

## Queue managers

Queues are managed by **queue managers**. An application can put messages only into a **local** queue. Getting a message is possible by extracting it from a **local** queue only  $\Rightarrow$  queue managers need to **route** messages.

## Routing



# Message broker


## Observation

Message queuing systems assume a **common messaging protocol**: all applications agree on message format (i.e., structure and data representation)

## Broker handles application heterogeneity in an MQ system

- Transforms incoming messages to target format
- Very often acts as an **application gateway**
- May provide **subject-based** routing capabilities (i.e., **publish-subscribe** capabilities)

# Message broker: general architecture






## Example: AMQP

### Lack of standardization

**Advanced Message-Queuing Protocol** was intended to play the same role as, for example, TCP in networks: a protocol for high-level messaging with different implementations.



### Basic model

Client sets up a (stable) **connection**, which is a container for several (possibly ephemeral) **one-way channels**. Two one-way channels can form a **session**. A **link** is akin to a socket, and maintains state about message transfers.

## Example: AMQP-based producer

```
1 import rabbitpy
2
3 def producer():
4     connection = rabbitpy.Connection() # Connect to RabbitMQ server
5     channel = connection.channel()    # Create new channel on the connection
6
7     exchange = rabbitpy.Exchange(channel, 'exchange') # Create an exchange
8     exchange.declare()
9
10    queue1 = rabbitpy.Queue(channel, 'example1') # Create 1st queue
11    queue1.declare()
12
13    queue2 = rabbitpy.Queue(channel, 'example2') # Create 2nd queue
14    queue2.declare()
15
16    queue1.bind(exchange, 'example-key') # Bind queue1 to a single key
17    queue2.bind(exchange, 'example-key') # Bind queue2 to the same key
18
19    message = rabbitpy.Message(channel, 'Test message')
20    message.publish(exchange, 'example-key') # Publish the message using the key
21    exchange.delete()
```

## Example: AMQP-based consumer

```
1 import rabbitpy
2
3 def consumer():
4     connection = rabbitpy.Connection()
5     channel = connection.channel()
6
7     queue = rabbitpy.Queue(channel, 'example1')
8
9     # While there are messages in the queue, fetch them using Basic.Get
10    while len(queue) > 0:
11        message = queue.get()
12        print('Message Q1: %s' % message.body.decode())
13        message.ack()
14
15    queue = rabbitpy.Queue(channel, 'example2')
16
17    while len(queue) > 0:
18        message = queue.get()
19        print('Message Q2: %s' % message.body.decode())
20        message.ack()
```

# Application-level multicasting

## Essence

Organize nodes of a distributed system into an **overlay network** and use that network to disseminate data:

- Oftentimes a **tree**, leading to unique paths
- Alternatively, also **mesh networks**, requiring a form of **routing**


# Application-level multicasting in Chord

## Basic approach

- ① Initiator generates a **multicast identifier**  $mid$ .
- ② Lookup  $succ(mid)$ , the node responsible for  $mid$ .
- ③ Request is routed to  $succ(mid)$ , which will become the **root**.
- ④ If  $P$  wants to join, it sends a **join** request to the root.
- ⑤ When request arrives at  $Q$ :
  - $Q$  has not seen a join request before  $\Rightarrow$  it becomes **forwarder**;  $P$  becomes child of  $Q$ . **Join request continues to be forwarded.**
  - $Q$  knows about tree  $\Rightarrow P$  becomes child of  $Q$ . **No need to forward join request anymore.**

# ALM: Some costs

## Different metrics




- **Link stress:** How often does an ALM message cross the same physical link? **Example:** message from A to D needs to cross  $\langle Ra, Rb \rangle$  twice.
- **Stretch:** Ratio in delay between ALM-level path and network-level path. **Example:** messages B to C follow path of length 73 at ALM, but 47 at network level  $\Rightarrow$  stretch =  $73/47$ .

# Flooding

## Essence


$P$  simply sends a message  $m$  to each of its neighbors. Each neighbor will forward that message, except to  $P$ , and only if it had not seen  $m$  before.



# Flooding

## Essence

$P$  simply sends a message  $m$  to each of its neighbors. Each neighbor will forward that message, except to  $P$ , and only if it had not seen  $m$  before.



## Variation

Let  $Q$  forward a message with a certain probability  $p_{flood}$ , possibly even dependent on its own number of neighbors (i.e., **node degree**) or the degree of its neighbors.



# Epidemic protocols

## Assume there are no write–write conflicts

- Update operations are performed at a single server
- A replica passes updated state to only a few neighbors
- Update propagation is lazy, i.e., not immediate
- Eventually, each update should reach every replica

## Two forms of epidemics

- **Anti-entropy**: Each replica regularly chooses another replica at random, and exchanges state differences, leading to identical states at both afterwards
- **Rumor spreading**: A replica which has just been updated (i.e., has been **contaminated**), tells several other replicas about its update (contaminating them as well).

# Anti-entropy

## Principle operations

- A node  $P$  selects another node  $Q$  from the system at random.
- **Pull**:  $P$  only pulls in new updates from  $Q$
- **Push**:  $P$  only pushes its own updates to  $Q$
- **Push-pull**:  $P$  and  $Q$  send updates to each other

## Observation

For push-pull it takes  $\mathcal{O}(\log(N))$  rounds to disseminate updates to all  $N$  nodes (**round** = when every node has taken the initiative to start an exchange).

# Anti-entropy: analysis

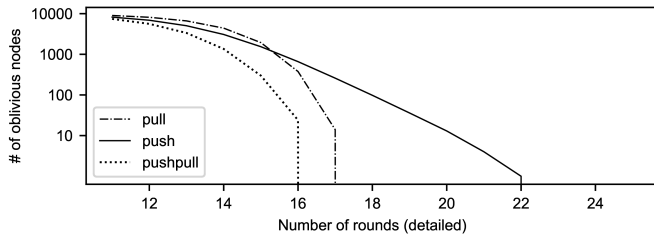
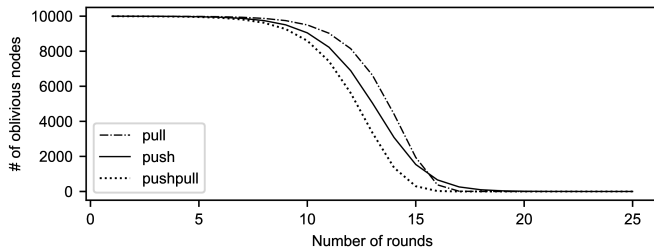
## Basics

Consider a single source, propagating its update. Let  $p_i$  be the probability that a node has not received the update after the  $i^{th}$  round.

## Analysis: staying ignorant

- With **pull**,  $p_{i+1} = (p_i)^2$ : the node was not updated during the  $i^{th}$  round and should contact another ignorant node during the next round.
- With **push**,  $p_{i+1} = p_i(1 - \frac{1}{N-1})^{(N-1)(1-p_i)} \approx p_i e^{-1}$  (for small  $p_i$  and large  $N$ ): the node was ignorant during the  $i^{th}$  round and no updated node chooses to contact it during the next round.
- With **push-pull**:  $(p_i)^2 \cdot (p_i e^{-1})$

# Anti-entropy performance



# Rumor spreading

## Basic model

A server  $S$  having an update to report, contacts other servers. If a server is contacted to which the update has already propagated,  $S$  stops contacting other servers with probability  $p_{stop}$ .

## Observation

If  $s$  is the fraction of ignorant servers (i.e., which are unaware of the update), it can be shown that with many servers

$$s = e^{-(1/p_{stop}+1)(1-s)}$$

# Formal analysis

## Notations

Let  $s$  denote fraction of nodes that have not yet been updated (i.e., **susceptible**;  $i$  the fraction of updated (**infected**) and active nodes; and  $r$  the fraction of updated nodes that gave up (**removed**).

## From theory of epidemics

$$(1) \quad ds/dt = -s \cdot i$$

$$(2) \quad di/dt = s \cdot i - p_{stop} \cdot (1 - s) \cdot i$$

$$\Rightarrow \quad di/ds = -(1 + p_{stop}) + \frac{p_{stop}}{s}$$

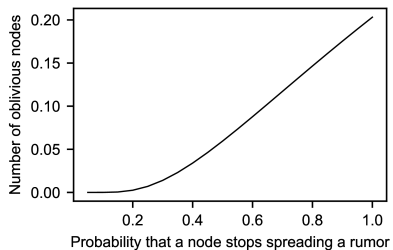
$$\Rightarrow \quad i(s) = -(1 + p_{stop}) \cdot s + p_{stop} \cdot \ln(s) + C$$

## Wrap up

$i(1) = 0 \Rightarrow C = 1 + p_{stop} \Rightarrow i(s) = (1 + p_{stop}) \cdot (1 - s) + p_{stop} \cdot \ln(s)$ . We are looking for the case  $i(s) = 0$ , which leads to  $s = e^{-(1/p_{stop}+1)(1-s)}$

# Rumor spreading


## The effect of stopping



Consider 10,000 nodes		
$1/p_{stop}$	$s$	$N_s$
1	0.203188	2032
2	0.059520	595
3	0.019827	198
4	0.006977	70
5	0.002516	25
6	0.000918	9
7	0.000336	3

# Rumor spreading

## The effect of stopping



Consider 10,000 nodes		
$1/p_{stop}$	$s$	$N_s$
1	0.203188	2032
2	0.059520	595
3	0.019827	198
4	0.006977	70
5	0.002516	25
6	0.000918	9
7	0.000336	3

## Note

If we really have to ensure that all servers are eventually updated, rumor spreading alone is not enough



# Deleting values

## Fundamental problem

We cannot remove an old value from a server and expect the removal to propagate. Instead, mere removal will be undone in due time using epidemic algorithms

## Solution

Removal has to be registered as a special update by inserting a **death certificate**

# Deleting values

## When to remove a death certificate (it is not allowed to stay for ever)

- Run a global algorithm to detect whether the removal is known everywhere, and then collect the death certificates (looks like garbage collection)
- Assume death certificates propagate in finite time, and associate a maximum lifetime for a certificate (can be done at risk of not reaching all servers)

## Note

It is necessary that a removal actually reaches all servers.