

# Principles of Distributed Systems

inft-3507

Dr. J.Burns

**ADA University**

Autumn 2025

## **Section 5: Coordination**

Ref: <https://www.distributed-systems.net/index.php/books/ds4/>

# Mutual exclusion

## Problem

Several processes in a distributed system want exclusive access to some resource.


## Basic solutions

**Permission-based:** A process wanting to enter its critical region, or access a resource, needs permission from other processes.


**Token-based:** A token is passed between processes. The one who has the token may proceed in its critical region, or pass it on when not interested.

# Permission-based, centralized


## Simply use a coordinator



(a)



(b)



(c)

- (a) Process  $P_1$  asks the coordinator for permission to access a shared resource. Permission is granted.
- (b) Process  $P_2$  then asks permission to access the same resource. The coordinator does not reply.
- (c) When  $P_1$  releases the resource, it tells the coordinator, which then replies to  $P_2$ .

## Mutual exclusion: Ricart & Agrawala

The same as Lamport except that acknowledgments are not sent


Return a response to a request only when:

- The receiving process has no interest in the shared resource; or
- The receiving process is waiting for the resource, but has lower priority (known through comparison of timestamps).


In all other cases, reply is *deferred*, implying some more local administration.

# Mutual exclusion: Ricart & Agrawala


## Example with three processes



(a)



(b)



(c)


- (a) Two processes want to access a shared resource at the same moment.
- (b)  $P_0$  has the lowest timestamp, so it wins.
- (c) When process  $P_0$  is done, it sends an *OK* also, so  $P_2$  can now go ahead.

# Mutual exclusion: Token ring algorithm

## Essence

Organize processes in a **logical** ring, and let a token be passed between them. The one that holds the token is allowed to enter the critical region (if it wants to).

An overlay network constructed as a logical ring with a circulating token



# Decentralized mutual exclusion

## Principle

Assume every resource is replicated  $N$  times, with each replica having its own coordinator  $\Rightarrow$  access requires a **majority vote** from  $m > N/2$  coordinators. A coordinator always responds immediately to a request.

## Assumption

When a coordinator crashes, it will recover quickly, but will have forgotten about permissions it had granted.

# Decentralized mutual exclusion

## How robust is this system?

- Let  $p = \Delta t / T$  be the probability that a coordinator resets during a time interval  $\Delta t$ , while having a lifetime of  $T$ .
- The probability  $\mathbb{P}[k]$  that  $k$  out of  $m$  coordinators reset during the same interval is

$$\mathbb{P}[k] = \binom{m}{k} p^k (1-p)^{m-k}$$

- $f$  coordinators reset  $\Rightarrow$  correctness is violated when there is only a minority of nonfaulty coordinators: when  $N - (m - f) \geq m$ , or,  $f \geq 2m - N$ .
- The probability of a violation is  $\sum_{k=2m-N}^m \mathbb{P}[k]$ .



# Decentralized mutual exclusion

## Violation probabilities for various parameter values

N	m	p	Violation
8	5	3 sec/hour	$< 10^{-5}$
8	6	3 sec/hour	$< 10^{-11}$
16	9	3 sec/hour	$< 10^{-4}$
16	12	3 sec/hour	$< 10^{-21}$
32	17	3 sec/hour	$< 10^{-4}$
32	24	3 sec/hour	$< 10^{-43}$

N	m	p	Violation
8	5	30 sec/hour	$< 10^{-3}$
8	6	30 sec/hour	$< 10^{-7}$
16	9	30 sec/hour	$< 10^{-2}$
16	12	30 sec/hour	$< 10^{-13}$
32	17	30 sec/hour	$< 10^{-2}$
32	24	30 sec/hour	$< 10^{-27}$

So....

What can we conclude?

# Mutual exclusion: comparison

Algorithm	Messages per entry/exit	Delay before entry (in message times)
Centralized	3	2
Distributed	$2(N - 1)$	$2(N - 1)$
Token ring	$1, \dots, \infty$	$0, \dots, N - 1$
Decentralized	$2kN + (k - 1)N/2 + N, k = 1, 2, \dots$	$2kN + (k - 1)N/2$

# Example: ZooKeeper

## Basics (and keeping it simple)

- Centralized server setup
- All client-server communication is **nonblocking**: a client immediately gets a response
- ZooKeeper maintains a **tree-based namespace**, akin to that of a filesystem
- Clients can **create**, **delete**, or **update** nodes, as well as **check existence**.

# ZooKeeper race condition

## Note

ZooKeeper allows a client to be **notified** when a node, or a branch in the tree, changes. This may easily lead to **race conditions**.


## Consider a simple locking mechanism

- 1 A client  $C_1$  creates a node `/lock`.
- 2 A client  $C_2$  wants to acquire the lock but is notified that the associated node already exists.
- 3 Before  $C_2$  subscribes to a notification,  $C_1$  releases the lock, i.e., deletes `/lock`.
- 4 Client  $C_2$  subscribes to changes to `/lock` and blocks locally.

## Solution

Use version numbers

# ZooKeeper versioning



## Notations

- $W(n,k)a$ : request to write  $a$  to node  $n$ , assuming current version is  $k$ .
- $R(n,k)$ : current version of node  $n$  is  $k$ .
- $R(n)$ : client wants to know the current value of node  $n$
- $R(n,k)a$ : value  $a$  from node  $n$  is returned with its current version  $k$ .

# ZooKeeper locking protocol

It is now very simple

- 1 **lock**: A client  $C_1$  creates a node `/lock`.
- 2 **lock**: A client  $C_2$  wants to acquire the lock but is notified that the associated node already exists  
 $\Rightarrow C_2$  subscribes to notification on changes of `/lock`.
- 3 **unlock**: Client  $C_1$  deletes node `/lock`  $\Rightarrow$  all subscribers to changes are notified.

# Election algorithms

## Principle

An algorithm requires that some process acts as a coordinator. The question is how to select this special process **dynamically**.

## Note

In many systems, the coordinator is chosen manually (e.g., file servers). This leads to centralized solutions  $\Rightarrow$  single point of failure.

## Teasers

- 1 If a coordinator is chosen dynamically, to what extent can we speak about a centralized or distributed solution?
- 2 Is a fully distributed solution, i.e. one without a coordinator, always more robust than any centralized/coordinated solution?

# Basic assumptions

- All processes have unique id's
- All processes know id's of all processes in the system (but not if they are up or down)
- Election means identifying the process with the highest id that is up



# Election by bullying


## Principle

Consider  $N$  processes  $\{P_0, \dots, P_{N-1}\}$  and let  $id(P_k) = k$ . When a process  $P_k$  notices that the coordinator is no longer responding to requests, it initiates an election:

- 1  $P_k$  sends an *ELECTION* message to all processes with higher identifiers:  $P_{k+1}, P_{k+2}, \dots, P_{N-1}$ .
- 2 If no one responds,  $P_k$  wins the election and becomes coordinator.
- 3 If one of the higher-ups answers, it takes over and  $P_k$ 's job is done.

# Election by bullying

## The bully election algorithm



# Election in a ring


## Principle

Process priority is obtained by organizing processes into a (logical) ring. The process with the highest priority should be elected as coordinator.

- Any process can start an election by sending an election message to its successor. If a successor is down, the message is passed on to the next successor.
- If a message is passed on, the sender adds itself to the list. When it gets back to the initiator, everyone had a chance to make its presence known.
- The initiator sends a coordinator message around the ring containing a list of all living processes. The one with the highest priority is elected as coordinator.

# Election in a ring

## Election algorithm using a ring



- The solid line shows the election messages initiated by  $P_6$
- The dashed one, the messages by  $P_3$

# Example: Leader election in ZooKeeper server group

## Basics

- Each server  $s$  in the server group has an identifier  $id(s)$
- Each server has a monotonically increasing counter  $tx(s)$  of the latest transaction it handled (i.e., series of operations on the namespace).
- When follower  $s$  suspects leader crashed, it broadcasts an *ELECTION* message, along with the pair  $(voteID, voteTX)$ . Initially,
  - $voteID \leftarrow id(s)$
  - $voteTX \leftarrow tx(s)$
- Each server  $s$  maintains two variables:
  - $leader(s)$ : records the server that  $s$  believes may be final leader. Initially,  $leader(s) \leftarrow id(s)$ .
  - $lastTX(s)$ : what  $s$  knows to be the most recent transaction. Initially,  $lastTX(s) \leftarrow tx(s)$ .

## Example: Leader election in ZooKeeper server group

When  $s^*$  receives  $(voteID, voteTX)$

- If  $lastTX(s^*) < voteTX$ , then  $s^*$  just received more up-to-date information on the most recent transaction, and sets
  - $leader(s^*) \leftarrow voteID$
  - $lastTX(s^*) \leftarrow voteTX$
- If  $lastTX(s^*) = voteTX$  and  $leader(s^*) < voteID$ , then  $s^*$  knows as much about the most recent transaction as what it was just sent, but its perspective on which server will be the next leader needs to be updated:
  - $leader(s^*) \leftarrow voteID$

### Note

When  $s^*$  believes it should be the leader, it broadcasts  $\langle id(s^*), tx(s^*) \rangle$ . Essentially, we're bullying.

# Example: Leader election in Raft

## Basics

- We have a (relatively small) group of servers
- A server is in one of three states: *follower*, *candidate*, or *leader*
- The protocol works in *terms*, starting with term 0
- Each server starts in the *follower* state.
- A leader is to regularly broadcast messages (perhaps just a simple *heartbeat*)

## Example: Leader election in Raft

### Selecting a new leader

When follower  $s^*$  hasn't received anything from the alleged leader  $s$  for some time,  $s^*$  broadcasts that it volunteers to be the next leader, increasing the term by 1.  $s^*$  enters the **candidate** state. Then:

- If leader  $s$  receives the message, it responds by acknowledging that it is still the leader.  $s^*$  returns to the **follower** state.
- If another follower  $s^{**}$  gets the election message from  $s^*$ , and it is the first election message during the current term,  $s^{**}$  votes for  $s^*$ . Otherwise, it simply ignores the election message from  $s^*$ .  
When  $s^*$  has collected a majority of votes, a new term starts with a new leader.

### Observation

By slightly differing the timeout values per follower for deciding when to start an election, we can avoid concurrent elections, and the election will rapidly converge.



# Elections by proof of work

## Basics

- Consider a potentially large group of processes
- Each process is required to solve a computational puzzle
- When a process solves the puzzle, it broadcasts its victory to the group
- We assume there is a conflict resolution procedure when more than one process claims victory

## Solving a computational puzzle

- Make use of a **secure hashing function**  $H(m)$ :
  - $m$  is some data;  $H(m)$  returns a **fixed-length bit string**
  - computing  $h = H(m)$  is computationally efficient
  - finding a function  $H^{-1}$  such that  $m = H^{-1}(H(m))$  is computationally extremely difficult
- **Practice**: finding  $H^{-1}$  boils down to an extensive **trial-and-error** procedure

# Elections by proof of work

## Controlled race

- Assume a globally known secure hash function  $H^*$ . Let  $H_i$  be the hash function used by process  $P_i$ .
- Task: given a bit string  $h = H_i(m)$ , find a bit string  $\tilde{h}$  such that  $h^* = H^*(H_i(\tilde{h} \odot h))$  where:
  - $h^*$  is a bit string with  $K$  leading zeroes
  - $\tilde{h} \odot h$  denotes some predetermined bitwise operation on  $\tilde{h}$  and  $h$

## Observation

By controlling  $K$ , we control the difficulty of finding  $\tilde{h}$ . If  $p$  is the probability that a random guess for  $\tilde{h}$  will suffice:  $p = (1/2)^K$ .

## Current practice

In many PoW-based blockchain systems,  $K = 64$

- With  $K = 64$ , it takes about 10 minutes on a supercomputer to find  $\tilde{h}$
- With  $K = 64$ , it takes about 100 years on a laptop to find  $\tilde{h}$

# Elections by proof of stake

## Basics

We assume a blockchain system in which  $N$  **secure tokens** are used:

- Each token has a unique **owner**
- Each token has a uniquely associated **index**  $1 \leq k \leq N$
- A token cannot be modified or copied without this going unnoticed

## Principle



- Draw a random number  $k \in \{1, \dots, N\}$
- Look up the process  $P$  that owns the token with index  $k$ .  $P$  is the next leader.

## Observation

The more tokens a process owns, the higher the probability it will be selected as leader.

# A solution for wireless networks

## A sample network





## Essence

Find the node with the highest capacity to select as the next leader.



# A solution for wireless networks

## A sample network



# A solution for wireless networks

## A sample network



## Essence

A node reports back only the node that it found to have the highest capacity.

# Gossip-based coordination: aggregation

## Typical apps

- **Data dissemination**: Perhaps the most important one. Note that there are many variants of dissemination.
- **Aggregation Function (Average)**: Let every node  $P_i$  maintain a variable  $v_i$ . When two nodes gossip, they each reset their variable to

$$v_i, v_j \leftarrow (v_i + v_j)/2$$

Result: in the end each node will have computed the average  $\bar{v} = \sum_i v_i / N$ .

- What happens in the case that initially  $v_i = 1$  and  $v_j = 0, j \neq i$ ?

# Gossip-based coordination: peer sampling

## Problem

For many gossip-based applications, you need to **select a peer uniformly at random** from the entire network. In principle, this means you need to know all other peers. **Impossible?**

## Basics

- Each node maintains a list of  $c$  references to other nodes
- **Regularly**, pick another node at random (from the list), and **exchange** roughly  $c/2$  references
- When the **application** needs to select a node at random, it also picks a random one from from its local list.

## Observation


Statistically, it turns out that the selection of a peer from the local list is indistinguishable from selecting uniformly at random peer from the entire network



# Gossip-based overlay construction

## Essence


Maintain two local lists of neighbors. The lowest is used for providing a **peer-sampling service**; the highest list is used to carefully select **application-dependent neighbors**.




# Gossip-based overlay construction: a 2D torus

Consider a logical  $N \times N$  grid, with a node on each point of the grid.


- Every node must maintain a list of  $c$  nearest neighbors
- Distance between node at  $(a_1, a_2)$  and  $(b_1, b_2)$  is  $d_1 + d_2$ , with  $d_i = \min(N - |a_i - b_i|, |a_i - b_i|)$
- Every node picks a random other node from its lowest-level list, and keeps only the closest one in its top-level list.
- Once every node has picked and selected a random node, we move to the next **round**



start ( $N = 50$ )



after 5 rounds



after 20 rounds


# A gossip-based 2D torus in Python (outline)

```
1 def maintainViews():
2     for viewType in [viewOverlay, viewPSS]: # For each view, do the same
3         peer[viewType] = None
4         if time to maintain viewType: # This viewType needs to be updated
5             peer[viewType] = selectPeer(viewType) # Select a peer
6             links = selectLinks(viewType, peer[viewType]) # Select links
7             sendTo(peer[viewType], Request[viewType], links) # Send links asynchronously
8
9     while True:
10        block = (peer[viewOverlay] != None) or (peer[viewPSS] != None)
11        sender, msgType, msgData = recvFromAny(block) # Block if expecting something
12
13        if msg == None: # All work has been done, simply return from the call
14            return
15
16        for viewType in [viewOverlay, viewPSS]: # For each view, do the same
17            if msgType == Response[viewType]: # Response to previously sent links
18                updateOwnView(viewType, msgData) # Just update the own view
19
20            elif msgType == Request[viewType]: # Request for exchanging links
21                if peer[viewType] == None: # No outstanding exchange request
22                    links = selectLinks(viewType, sender) # Select links
23                    sendTo(sender, Response[viewType], links) # Send them asynchronously
24                    updateOwnView(viewType, msgData) # Update own view
25                else: # This node already has a pending exchange request, ignore this one
26                    sendTo(sender, IgnoreRequest[viewType])
27
28            elif msgType == IgnoreRequest[viewType]: # Request has been denied, give up
29                peer[viewType] = None
```

# Secure gossiping

## Dramatic attack

Consider when exchanging references, a set of **colluding nodes** systematically returns links only to each other  $\Rightarrow$  we are dealing with **hub attack**.






## Situation

A network with 100,000 nodes, a local list size  $c = 30$ , and only 30 attackers. The y-axis shows the number of nodes with links **only** to the attackers. After less than 300 rounds, the attackers have full control.

# A solution: gathering statistics

This is what measuring indegree distributions tells us: which fraction of nodes (y-axis) have how many other nodes pointing to them (x-axis)?




## Basic approach

When a benign node initiates an exchange, it may either use the result for gathering statistics, or for updating its local list. An attacker is in limbo: will its response be used for statistical purposes or for functional purposes?

## Observation

When gathering statistics may reveal colluders, a colluding node will be **forced** to behave according to the protocol.

# Distributed event matching



## Principle

- A process specifies in which events it is interested (**subscription  $S$** )
- When a process **publishes a notification  $N$**  we need to see whether  $S$  **matches  $N$** .

## Hard part

Implementing the **match** function in a scalable manner.

# General approach

## What is needed


- *sub2node*( $S$ ): map a subscription  $S$  to a nonempty subset  $\mathbf{S}$  of servers
- *not2node*( $N$ ): map a notification  $N$  to a nonempty subset  $\mathbf{N}$  of servers

Make sure that  $\mathbf{S} \cap \mathbf{N} \neq \emptyset$ .


## Observations

- Centralized solution is simple:  $\mathbf{S} = \mathbf{N} = \{s\}$ , i.e. a single server.
- Topic-based publish-subscribe is also simple: each  $S$  and  $N$  is tagged with a **single topic**; each topic is handled by a single server (a **rendezvous node**). Several topics may be handled by same server).
- Content-based publish-subscribe is **tough**: a subscription takes the form (*attribute*, *value*) pair, with example values:
  - **range**: " $1 \leq x < 10$ "
  - **containment**: " $x \in \{\text{red}, \text{blue}\}$ "
  - **prefix and suffix expressions**: "`url.startsWith("https")`"

# Selective routing



(a)



(b)

- (a) first broadcast subscriptions
- (b) forward notifications only to relevant rendezvous nodes

## Example of a (partially filled) routing table

Interface	Filter
To node 3	$a \in [0, 3]$
To node 4	$a \in [2, 5]$
Toward router $R_1$	(unspecified)



# Gossiping: Sub-2-Sub

## Basics

- **Goal:** To realize scalability, make sure that subscribers with the same interests form just a single group
- **Model:** There are  $N$  attributes  $a_1, \dots, a_N$ . An attribute value is always (mappable to) a floating-point number.
- **Subscription:** Takes forms such as  $S = \langle a_1 \rightarrow 3.0, a_4 \rightarrow [0.0, 0.5] \rangle$ :  $a_1$  should be 3.0;  $a_4$  should lie between 0.0 and 0.5; other attribute values don't matter.

## Observations


- A subscription  $S_i$  specifies a subset  $\mathbf{S}_i$  in a  $N$ -dimensional space.
- We are interested only in notifications that fall into  $\bar{\mathbf{S}} = \cup \mathbf{S}_i$ .

## Goal

Partition  $\bar{\mathbf{S}}$  into  $M$  disjoint subspaces  $\bar{\mathbf{S}}_1, \dots, \bar{\mathbf{S}}_M$  such that

- **Partitioning:**  $\forall k \neq m : \bar{\mathbf{S}}_k \cap \bar{\mathbf{S}}_m = \emptyset$  and  $\cup_m \bar{\mathbf{S}}_m = \bar{\mathbf{S}}$
- **Subscription coverage:**  $(\bar{\mathbf{S}}_m \cap \mathbf{S}_i \neq \emptyset) \Rightarrow (\bar{\mathbf{S}}_m \subseteq \mathbf{S}_i)$

# Gossiping: Sub-2-Sub



## Consider a single attribute

- Nodes regularly exchange their subscriptions through gossiping
- An intersection between two nodes leads to a mutual reference
- If  $S_{ijk} = S_i \cap S_j \cap S_k \neq \emptyset$  and  $S_{ij} - S_{ijk} \neq \emptyset$ , then:
  - nodes  $i, j, k$  are grouped into a **single overlay network** (for  $S_{ijk}$ )
  - nodes  $i, j$  are grouped into a **single overlay network** (for  $S_{ij} - S_{ijk}$ )

# Secure publish-subscribe

## We are facing nasty dilemma's

- **Referential decoupling**: messages should be able to flow from a publisher to subscribers while guaranteeing mutual anonymity  $\Rightarrow$  we cannot set up a secure channel.
- Not knowing where messages come from imposes **integrity problems**.
- Assuming a **trusted broker** may easily be practically impossible, certainly when dealing with sensitive information  $\Rightarrow$  we now have a **routing problem**.

## Solution

- Allow for searching (and matching) on encrypted data, without the need for decryption.
- **PEKS**: accompany encrypted messages with a collection of (again encrypted) keywords and search for matches on keywords.


# Public-Key Encryption with Keyword Search (PEKS)

## Basics

- Use a public key  $PK$ , message  $m$  and its  $n$  keywords  $KW_1, \dots, KW_n$  are stored at a server as the message  $m^*$ :

$$m^* = [PK(m) | PEKS(PK, KW_1) | PEKS(PK, KW_2) | \dots | PEKS(PK, KW_n)]$$

- A subscriber gets the accompanying secret key.
- For each keyword  $KW_i$ , a **trapdoor**  $T_{KW_i}$  is generated:  $T_W(m^*)$  will return *true* iff  $W \in \{KW_1, \dots, KW_n\}$ .



# Positioning nodes

## Issue


In large-scale distributed systems in which nodes are dispersed across a wide-area network, we often need to take some notion of **proximity** or **distance** into account  $\Rightarrow$  it starts with determining a (relative) **location** of a node.

# Computing position

## Observation

A node  $P$  needs  $d + 1$  **landmarks** to compute its own position in a  $d$ -dimensional space. Consider two-dimensional case.

## Computing a position in 2D



## Solution

$P$  needs to solve three equations in two unknowns  $(x_P, y_P)$ :

$$d_i = \sqrt{(x_i - x_P)^2 + (y_i - y_P)^2}$$

# Global Positioning System

Assuming that the clocks of the satellites are accurate and synchronized

- It takes a while before a signal reaches the receiver
- The receiver's clock is definitely out of sync with the satellite

## Basics

- $\Delta_r$ : **unknown deviation** of the receiver's clock.
- $x_r, y_r, z_r$ : **unknown coordinates** of the receiver.
- $T_i$ : timestamp on a message from satellite  $i$
- $\Delta_i = (T_{now} - T_i) + \Delta_r$ : **measured delay** of the message sent by satellite  $i$ .
- **Measured distance** to satellite  $i$ :  $c \times \Delta_i$  ( $c$  is speed of light)
- Real distance:  $d_i = c\Delta_i - c\Delta_r = \sqrt{(x_i - x_r)^2 + (y_i - y_r)^2 + (z_i - z_r)^2}$

## Observation

4 satellites  $\Rightarrow$  4 equations in 4 unknowns (with  $\Delta_r$  as one of them)

# WiFi-based location services

## Basic idea

- Assume we have a database of known access points (APs) with coordinates
- Assume we can estimate distance to an AP
- Then: with 3 detected access points, we can compute a position.

## War driving: locating access points

- Use a WiFi-enabled device along with a GPS receiver, and move through an area while recording observed access points.
- Compute the centroid: assume an access point  $AP$  has been detected at  $N$  different locations  $\{\vec{x}_1, \vec{x}_2, \dots, \vec{x}_N\}$ , with known GPS location.
- Compute location of  $AP$  as  $\vec{x}_{AP} = \frac{\sum_{i=1}^N \vec{x}_i}{N}$ .

## Problems

- Limited accuracy of each GPS detection point  $\vec{x}_i$
- An access point has a nonuniform transmission range
- Number of sampled detection points  $N$  may be too low.




# Computing position

## Problems

- Measured latencies to landmarks fluctuate
- Computed distances will not even be consistent

## Inconsistent distances in 1D space



## Solution: minimize errors

- Use  $N$  special **landmark nodes**  $L_1, \dots, L_N$ .
- Landmarks measure their pairwise latencies  $\tilde{d}(L_i, L_j)$
- A central node computes the coordinates for each landmark, minimizing:

$$\sum_{i=1}^N \sum_{j=i+1}^N \left( \frac{\tilde{d}(L_i, L_j) - \hat{d}(L_i, L_j)}{\tilde{d}(L_i, L_j)} \right)^2$$

where  $\hat{d}(L_i, L_j)$  is distance after nodes  $L_i$  and  $L_j$  have been positioned

# Computing position

## Choosing the dimension $m$

The hidden parameter is the dimension  $m$  with  $N > m$ . A node  $P$  measures its distance to each of the  $N$  landmarks and computes its coordinates by minimizing

$$\sum_{i=1}^N \left( \frac{\tilde{d}(L_i, P) - \hat{d}(L_i, P)}{\tilde{d}(L_i, P)} \right)^2$$

## Observation

Practice shows that  $m$  can be as small as 6 or 7 to achieve latency estimations within a factor 2 of the actual value.

# Vivaldi

## Principle: network of springs exerting forces

Consider a collection of  $N$  nodes  $P_1, \dots, P_N$ , each  $P_i$  having coordinates  $\vec{x}_i$ . Two nodes exert a **mutual force**:

$$\vec{F}_{ij} = (\tilde{d}(P_i, P_j) - \hat{d}(P_i, P_j)) \times u(\vec{x}_i - \vec{x}_j)$$

with  $u(\vec{x}_i - \vec{x}_j)$  is the unit vector in the direction of  $\vec{x}_i - \vec{x}_j$

## Node $P_i$ repeatedly executes steps

- 1 Measure the latency  $\tilde{d}_{ij}$  to node  $P_j$ , and also receive  $P_j$ 's coordinates  $\vec{x}_j$ .
- 2 Compute the error  $e = \tilde{d}(P_i, P_j) - \hat{d}(P_i, P_j)$
- 3 Compute the direction  $\vec{u} = u(\vec{x}_i - \vec{x}_j)$ .
- 4 Compute the force vector  $F_{ij} = e \cdot \vec{u}$
- 5 Adjust own position by moving along the force vector:  $\vec{x}_i \leftarrow \vec{x}_i + \delta \cdot \vec{u}$ .