

# Principles of Distributed Systems

inft-3507

Dr. J.Burns

**ADA University**

Autumn 2025

## **Section 6: Naming**

Ref: <https://www.distributed-systems.net/index.php/books/ds4/>

# Names, identifiers, and addresses

# Naming

## Essence

Names are used to denote entities in a distributed system. To operate on an entity, we need to access it at an **access point**. Access points are entities that are named by means of an **address**.

## Note

A **location-independent** name for an entity  $E$ , is independent of the addresses of the access points offered by  $E$ .

# Identifiers

## Pure name

A name that has no meaning at all; it is just a random string. Pure names can be used for comparison only.

## Identifier: A name having some specific properties

- 1 An identifier refers to at most one entity.
- 2 Each entity is referred to by at most one identifier.
- 3 An identifier always refers to the same entity (i.e., it is never reused).

## Observation

An identifier need not necessarily be a pure name, i.e., it may have content.

## Flat naming

# Broadcasting

Broadcast the ID, requesting the entity to return its current address

- Can never scale beyond local-area networks
- Requires all processes to listen to incoming location requests

## Address Resolution Protocol (ARP)

To find out which MAC address is associated with an IP address, broadcast the query “who has this IP address”?

# Forwarding pointers

When an entity moves, it leaves behind a pointer to its next location

- Dereferencing can be made entirely transparent to clients by simply following the chain of pointers
- Update a client's reference when present location is found
- Geographical scalability problems (for which separate chain reduction mechanisms are needed):
  - Long chains are not fault tolerant
  - Increased network latency at dereferencing

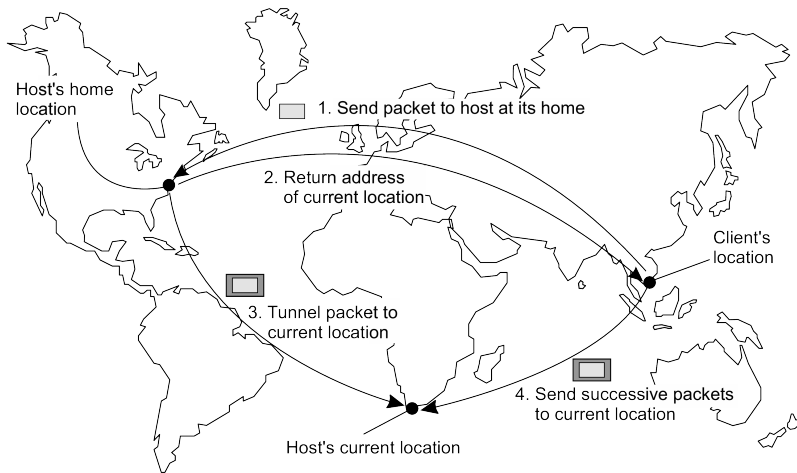
# Home-based approaches

Single-tiered scheme: Let a **home** keep track of where the entity is

- Entity's **home address** registered at a naming service
- The home registers the **foreign address** of the entity
- Client contacts the home first, and then continues with foreign location



# The principle of mobile IP



# Home-based approaches

## Problems with home-based approaches

- Home address has to be supported for entity's lifetime
- Home address is fixed  $\Rightarrow$  unnecessary burden when the entity permanently moves
- Poor geographical scalability (entity may be next to client)

## Note

Permanent moves may be tackled with another level of naming (DNS)

# Illustrative: Chord

Consider the organization of many nodes into a **logical ring**

- Each node is assigned a random  $m$ -bit **identifier**.
- Every entity is assigned a unique  $m$ -bit **key**.
- Entity with key  $k$  falls under jurisdiction of node with smallest  $id \geq k$  (called its **successor**  $succ(k)$ ).

## Nonsolution

Let each node keep track of its neighbor and start linear search along the ring.

## Notation

We will speak of node  $p$  as the node have identifier  $p$

# Chord finger tables

## Principle

- Each node  $p$  maintains a **finger table**  $FT_p[]$  with at most  $m$  entries:

$$FT_p[i] = \text{succ}(p + 2^{i-1})$$

**Note:** the  $i$ -th entry points to the first node succeeding  $p$  by at least  $2^{i-1}$ .


- To look up a key  $k$ , node  $p$  forwards the request to node with index  $j$  satisfying

$$q = FT_p[j] \leq k < FT_p[j+1]$$

- If  $p < k < FT_p[1]$ , the request is also forwarded to  $FT_p[1]$

# Chord lookup example

Resolving key 26 from node 1 and key 12 from node 28



# Exploiting network proximity

## Problem

The logical organization of nodes in the overlay may lead to **erratic message transfers** in the underlying Internet: node  $p$  and node  $\text{succ}(p+1)$  may be very far apart.

## Solutions


- **Topology-aware node assignment**: When assigning an ID to a node, make sure that nodes close in the ID space are also close in the network. **Can be very difficult.**
- **Proximity routing**: Maintain more than one possible successor, and forward to the closest.  
**Example**: in Chord  $FT_p[i]$  points to first node in  $INT = [p + 2^{i-1}, p + 2^i - 1]$ . Node  $p$  can also store pointers to other nodes in  $INT$ .
- **Proximity neighbor selection**: When there is a choice of selecting who your neighbor will be (not in Chord), pick the closest one.

# Hierarchical Location Services (HLS)

## Basic idea

Build a large-scale search tree for which the underlying network is divided into hierarchical domains. Each domain is represented by a separate directory node.

## Principle




# HLS: Tree organization

## Invariants

- Address of entity  $E$  is stored in a leaf or intermediate node
- Intermediate nodes contain a pointer to a child if and only if the subtree rooted at the child stores an address of the entity
- The root knows about all entities

## Storing information of an entity having two addresses in different leaf domains






# HLS: Lookup operation

## Basic principles

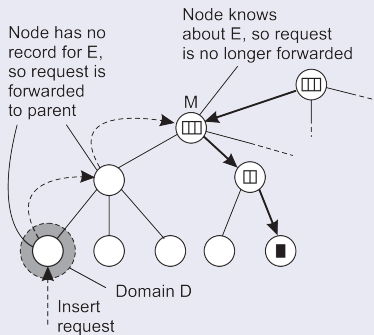
- Start lookup at local leaf node
- Node knows about  $E \Rightarrow$  follow downward pointer, else go up
- Upward lookup always stops at root

## Looking up a location




# HLS: Insert operation

(a) An insert request is forwarded to the first node that knows about entity  $E$ . (b) A chain of forwarding pointers to the leaf node is created



(a)



(b)

# Can an HLS scale?

## Observation

A design flaw seems to be that the root node needs to keep track of **all** identifiers  $\Rightarrow$  make a distinction between a **logical design** and its **physical implementation**.

## Notation

- Assume there are a total of  $N$  physical hosts  $\{H_1, H_2, \dots, H_N\}$ . Each host is capable of running one or more location servers.
- $D_k(A)$  denotes the domain at level  $k$  that contains address  $A$ ;  $k = 0$  denotes the root domain.
- $LS_k(E, A)$  denotes the unique location server in  $D_k(A)$  responsible for keeping track of entity  $E$ .

# Can an HLS scale?

## Basic idea for scaling


- Choose different physical servers for the logical name servers on a per-entity basis
  - (at root level, but also intermediate)
- Implement a mapping of entities to physical servers such that the load of storing records will be distributed

# Can an HLS scale?

## Solution

- $\mathbf{D}_k = \{D_{k,1}, D_{k,2}, \dots, D_{k,N_k}\}$  denotes the  $N_k$  domains at level  $k$
- **Note:**  $N_0 = |\mathbf{D}_0| = 1$ .
- For each level  $k$ , the set of hosts is partitioned into  $N_k$  subsets, with each host running a location server representing exactly one of the domains  $D_{k,i}$  from  $\mathbf{D}_k$ .

## Principle of distributing logical location servers



# Security in flat naming

## Basics

Without special measures, we need to trust that the name-resolution process to return what is associated with a flat name. Two approaches to follow:

- Secure the identifier-to-entity association
- Secure the name-resolution process

## Self-certifying names

Use a value derived from the associated entity and make it (part of) the flat name:

- $id(entity) = hash(data\ associated\ with\ the\ entity)$

when dealing with read-only entities, otherwise

- $id(entity) = public\ key(entity)$

in which case additional data is returned, such as a verifiable digital signature.

## Securing the name-resolution process

Much more involved: discussion deferred until discussing secure DNS


# Structured naming

# Name space

## Naming graph

A graph in which a **leaf node** represents a (named) entity. A **directory node** is an entity that refers to other nodes.

## A general naming graph with a single root node



## Note

A directory node contains a table of *(node identifier, edge label)* pairs.



# Name space

We can easily store all kinds of attributes in a node

- Type of the entity
- An identifier for that entity
- Address of the entity's location
- Nicknames
- ...

## Note

Directory nodes can also have attributes, besides just storing a directory table with *(identifier, label)* pairs.

# Name resolution

## Problem

To resolve a name, we need a directory node. How do we actually find that (initial) node?

Closure mechanism: The mechanism to select the implicit context from which to start name resolution

- `www.distributed-systems.net`: start at a DNS name server
- `/home/maarten/mbox`: start at the local NFS file server (possible recursive search)
- `0031 20 598 7784`: dial a phone number
- `77.167.55.6`: route message to a specific IP address

## Note

You cannot have an explicit closure mechanism – how would you start?

# Name linking

## Hard link

What we have described so far as a **path name**: a name that is resolved by following a specific path in a naming graph from one node to another.

Soft link: Allow a node  $N$  to contain a **name** of another node


- First resolve  $N$ 's name (leading to  $N$ )
- Read the content of  $N$ , yielding *name*
- Name resolution continues with *name*

## Observations

- The name resolution process determines that we read the **content** of a node, in particular, the name in the other node that we need to go to.
- One way or the other, we know where and how to start name resolution given *name*

# Name linking

## The concept of a symbolic link explained in a naming graph



### Observation

Node *n5* has only one name

# DNS


## Essence

- Hierarchically organized name space with each node having exactly one incoming edge  $\Rightarrow$  edge label = node label.
- **domain**: a subtree
- **domain name**: a path name to a domain's root node.


## Information in a node

Type	Refers to	Description
<i>SOA</i>	Zone	Holds info on the represented zone
<i>A</i>	Host	IP addr. of host this node represents
<i>MX</i>	Domain	Mail server to handle mail for this node
<i>SRV</i>	Domain	Server handling a specific service
<i>NS</i>	Zone	Name server for the represented zone
<i>CNAME</i>	Node	Symbolic link
<i>PTR</i>	Host	Canonical name of a host
<i>HINFO</i>	Host	Info on this host
<i>TXT</i>	Any kind	Any info considered useful

# Modern DNS



The traditional organization of the implementation of DNS



The modern organization of DNS

# Secure DNS

## Basic approach

Resource records of the same type are grouped into a **signed set**, per zone. Examples:


- A set with all the IPv4 addresses of a zone
- A set with all the IPv6 addresses of a zone
- A set with the name servers of a zone

The public key associated with the secret key used for signing a set of resource records is added to a zone, called a **zone-signing key**.

## Trusting the signatures

- All zone-signing keys are grouped again into a separate set, which is signed using another secret key. The public key of the latter is the **key-signing key**.
- The hash of the key-signing key **is stored at, and signed by, the parent zone**

# Secure DNS



## Building a trust chain


- Consider a **single set of resource records  $RR$** , hashed with  $HZ_k$  and signed with  $SKZ_k$
- $SKZ_k$  has associated public key  $ZSK_k$
- (Set of)  $ZSK_k$  is hashed with  $HK_k$  and signed with  $SKK_k$
- $SKK_k$  has associated public key  $KSK_k$

A client can **verify signature**  $SKZ_2(HZ_2(RR))$  by checking

$$ZSK_2(SKZ_2(HZ_2(RR))) \stackrel{?}{=} HZ_2(RR)$$



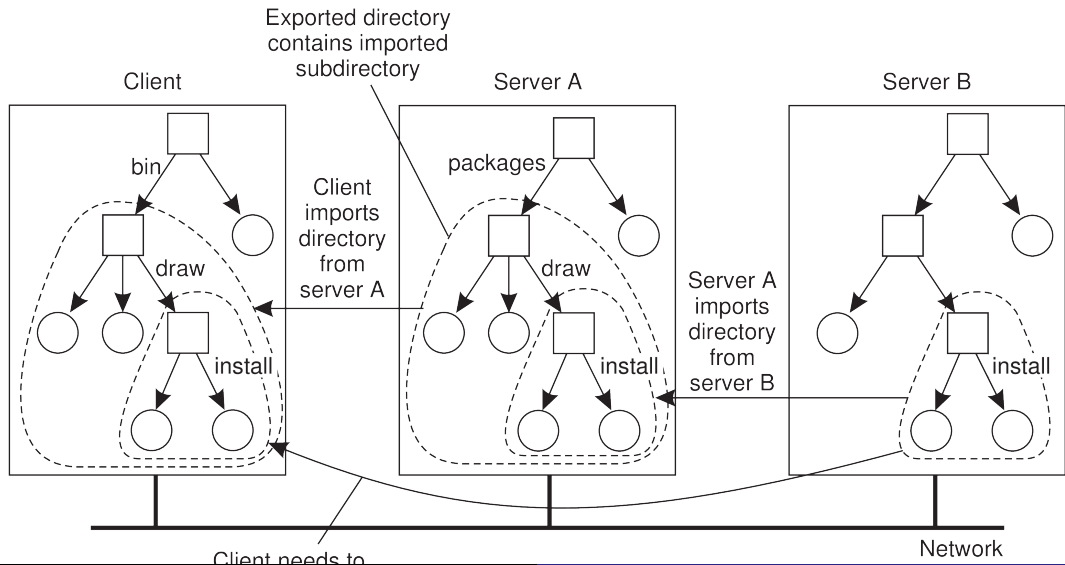
# Naming in NFS



## Observation

A server may **export** (a part of) its filesystem, which can then be **imported** by different clients by **mounting**. Note that different clients will have **different (nonsharable) namespaces**!

# Mounting nested directories



# Attribute-based naming

# Attribute-based naming

## Observation

In many cases, it is much more convenient to name, and look up entities through their **attributes**  $\Rightarrow$  traditional **directory services** (aka **yellow pages**).

## Problem

Lookup operations can be expensive, as they require matching **requested attribute values**, against **actual attribute values**  $\Rightarrow$  inspect **all entities** (in principle).

# Implementing directory services

## Solution for scalable searching

Implement basic directory service as database, and combine with traditional structured naming system.

## Lightweight Directory Access Protocol (LDAP)

Each directory entry consists of (*attribute, value*) pairs, and is **uniquely named** to ease lookups.


Attribute	Abbr.	Value
Country	<i>C</i>	NL
Locality	<i>L</i>	Amsterdam
Organization	<i>O</i>	VU University
OrganizationalUnit	<i>OU</i>	Computer Science
CommonName	<i>CN</i>	Main server
Mail_Servers	–	137.37.20.3, 130.37.24.6, 137.37.20.10
FTP_Server	–	130.37.20.20
WWW_Server	–	130.37.20.20

# LDAP

## Essence

- **Directory Information Base**: collection of all directory entries in an LDAP service.
- Each record is uniquely named as a sequence of naming attributes (called **Relative Distinguished Name**), so that it can be looked up.
- **Directory Information Tree**: the naming graph of an LDAP directory service; each node represents a directory entry.

## Part of a directory information tree



# LDAP

## Two directory entries having *HostName* as RDN

Attribute	Value	Attribute	Value
<i>Locality</i>	<i>Amsterdam</i>	<i>Locality</i>	<i>Amsterdam</i>
<i>Organization</i>	<i>VU University</i>	<i>Organization</i>	<i>VU University</i>
<i>OrganizationalUnit</i>	<i>Computer Science</i>	<i>OrganizationalUnit</i>	<i>Computer Science</i>
<i>CommonName</i>	<i>Main server</i>	<i>CommonName</i>	<i>Main server</i>
<i>HostName</i>	<i>star</i>	<i>HostName</i>	<i>zephyr</i>
<i>HostAddress</i>	<i>192.31.231.42</i>	<i>HostAddress</i>	<i>137.37.20.10</i>

Result of `search("(C=NL)(O=VU University)(OU=*)(CN=Main server)")`

# Distributed index

## Basic idea

- Assume a set of attributes  $\{a^1, \dots, a^N\}$
- Each attribute  $a^k$  takes values from a set  $R^k$
- For each attribute  $a^k$  associate a set  $\mathbf{S}^k = \{S_1^k, \dots, S_{n_k}^k\}$  of  $n_k$  servers
- **Global mapping**  $F$ :  $F(a^k, v) = S_j^k$  with  $S_j^k \in \mathbf{S}^k$  and  $v \in R^k$

## Observation

If  $L(a^k, v)$  is set of keys returned by  $F(a^k, v)$ , then a query can be formulated as a logical expression, e.g.,

$$(F(a^1, v^1) \wedge F(a^2, v^2)) \vee F(a^3, v^3)$$

which can be processed by the client by constructing the set

$$(L(a^1, v^1) \cap L(a^2, v^2)) \cup L(a^3, v^3)$$



# Drawbacks of distributed index

## Quite a few

- A query involving  $k$  attributes requires contacting  $k$  servers
- Imagine looking up "*lastName = Smith*  $\wedge$  *firstName = Pheriby*": the client may need to process **many** files as there are so many people named "Smith."
- No (easy) support for **range queries**, such as "*price* = [1000 – 2500]."

## Summary

# Summary

In the *Naming* section of the lecture notes we have discussed the following topics

- 1 Names, identifiers, and addresses
- 2 Flat Naming
- 3 Structured Naming
- 4 Attribute-based Naming